When Can Markets Crash Without News?

Siyuan Ye
Independent Researcher

October 2026

Abstract

In a resilient N-asset market with fixed fundamentals and paired-dollar execution, we prove that response elasticity $e(S) = Sg'(S)/g(S) \geq 2$ excludes every positive systemic drawdown, whereas $e(S) \geq 1$ excludes a simultaneous decline of all assets. The proof uses a response-adapted generalized price coordinate $\Phi(S) = \int_F^S u/g(u)\,\mathrm{d}u$, uniquely normalized so that an executed-dollar transfer produces an equal coordinate displacement. This reveals an additive trading conservation law and a recovery-driven stability barrier. Within the common family $g_p(S) = \eta(S/F)^p$, the regimes $p = 2, 1, 0$ respectively yield exclusion, inventory-bounded partial-coverage losses, and arbitrarily severe partial-coverage log losses under uniform finite turnover and inventory when pulses are permitted. We establish necessary dollar-turnover and share-inventory bounds. For constant per-share impact and bounded gross share rates, the maximal small-rate systemic loss is of quadratic order, with an attaining-order continuous construction. Rolling daily price limits constrain individual sessions but need not exclude multiday losses: a computer-assisted ten-asset path has nine assets reach their lower closing limits on two successive sessions, closes every executed-dollar ledger, and attains systemic log loss above 0.020. In this example the required gross-rate cap satisfies $0.044 < \bar{v}_* < 0.285$.

Keywords: market crashes; price impact; market resilience; systemic drawdown; reachability; execution constraints.
MSC 2020: 91G15, 91G45, 93B03, 65G30, 91G80.

1 Introduction

Large market declines need not follow bad news. Transactions themselves change marginal prices, and market resilience determines how much of that displacement survives. An elementary but consequential question is therefore whether an unchanged-fundamentals market can be forced into a *broad joint loss* by trades that obey its own settlement rules. We ask a narrower question than how an equilibrium chooses to crash: *which systemic loss states are reachable, and what executed trading resources are necessary to reach them?*

Existing accounts explain news-free price movements through distinct market mechanisms. Gennotte and Leland [9] analyze dynamic hedging and inference, while Romer [17] studies information revelation through trading without external news. Huang and Wang [10] show how nonsynchronous trading needs generate endogenous order imbalances; Cespa and Foucault [4] describe feedback in cross-asset liquidity. In a related multiasset setting, Cont and Wagalath [6] derive endogenous return dependence from distressed selling. These models establish why liquidity-driven instability matters. We study a complementary mathematical question: *which joint terminal losses are*

consistent with a stated execution law, recovery dynamics, and a zero terminal executed-dollar ledger for every asset?

The equilibrium price-response benchmark of Kyle [14] is one classical point of departure. Price-impact and recovery models provide the tools to make this question quantitative. Obizhaeva and Wang [16] and Alfonsi and Schied [2] characterize execution under resilient liquidity, and Gatheral [8], Alfonsi, Schied, and Slynko [3], and Huberman and Stanzl [11] derive restrictions from trading-cost and manipulation criteria. Multivariate transient impact is analyzed by Alfonsi, Klöck, and Schied [1], while Cordoni and Lillo [7] investigate equilibrium instability in multiagent, multiasset impact games. Kyle and Obizhaeva [15] relate large bets and execution speeds to observed crashes through market-microstructure invariance. We take these mechanisms as motivation for a different task: characterizing *feasible systemic loss states* under a specified settlement protocol, and then quantifying the turnover, inventory, rate, and price-band constraints on paths to those states.

We study a *cashless risky-asset reallocation pool*: each executed sale is paired with an equal-dollar risky-asset purchase, and every asset's terminal signed-dollar flow is zero. Prices recover toward a fixed benchmark F, while executed shares move marginal prices through a positive response function $g(S)$. The key step is to measure price in a *response-adapted coordinate*, in which the otherwise nonlinear effect of executed dollars becomes additive. Paired trading then conserves the sum of these coordinates; resilience changes that sum between trades. The coordinate's curvature yields exclusion barriers and unifies three response regimes under common constraints. We next ask which of the structurally permitted losses remain attainable once execution resources and exchange restrictions are imposed.

For bounded gross share rates and daily price bands, reachability requires more than a permitted price endpoint. When most assets fall, the others must accept the matched dollars, yet their prices and execution rates are also constrained; moreover their appreciation can offset the average loss entering the crisis criterion. We develop necessary rate and cross-subset capacity bounds and verify an explicit two-day path with nine out of ten assets repeatedly closing at a stylized lower limit. The result is a *certified multiday systemic drawdown witness* satisfying all stated path constraints. It provides an explicit existence counterpart to the necessary resource inequalities.

Our analysis proceeds from structural barriers to necessary resources, and then from execution restrictions to a simultaneous feasibility certificate. We first give the main statements; the common market definition and proofs follow.

Theorem 1.1 (Main results, informal)**.** *Under the paired-dollar, resilient market defined in Section 2, the following conclusions hold.*

(i) Structural exclusion. *If $e(S) = Sg'(S)/g(S) \geq 2$ at all admissible prices, then $\mathcal{C}_\alpha(t) \leq 0$ for every breadth and time. The weaker condition $e(S) \geq 1$ excludes a simultaneous decline of all assets.*

(ii) Reachability regimes. *With $g_p(S) = \eta(S/F)^p$, $p = 2, 1, 0$ respectively produce exclusion, inventory-bounded partial-coverage drawdowns, and, for $k = \lceil \alpha N \rceil < N$ and sufficiently long finite horizons, arbitrarily severe partial-coverage losses under fixed finite turnover and inventory caps when pulses are allowed. For constant impact and $k = N$, a sharp horizon-dependent bound replaces unrestricted severity.*

(iii) Execution resources. *General positive g yields explicit necessary dollar-turnover and share-inventory bounds. With constant per-share impact and bounded gross executed-share rates, the small-rate crisis bound is quadratic and this order is attained by a continuous dollar-closed construction.*

(iv) Multiday execution. *Rolling price bands restrict per-session losses, while rate limits and the receiving assets' capacity restrict consecutive lower-limit closes. A computer-assisted two-day ten-asset path satisfies these controls, has nine repeated lower-limit closes, and yields* $\mathcal{C}_{0.9}(T) > 0.020$*; for that instance,* $0.044 < \bar{v}_* < 0.285$.

The common mathematical tool is the *response-adapted generalized price coordinate*

$$\Phi(S) = \int_F^S \frac{u}{g(u)}\,\mathrm{d}u. \tag{1.1}$$

It is the coordinate in which an executed dollar produces one unit of price-state displacement. Its additive conservation during paired trades, together with a restoring drift between trades, connects the structural barriers to the resource and multiday execution results.

2 The market: assumptions, execution, and price dynamics

We specify a risky-asset reallocation submarket in which every executed sale is matched by an equal-dollar purchase. Settlement and price response are distinct assumptions: the former constrains flows, while the latter determines the induced price motion. All subsequent theorems use the same market definition.

Assumption 2.1 (The resilient paired-dollar market)**.** Fix $N \geq 2$ assets and a finite horizon $[0, T]$. The following jointly define the admissible market.

(M1) **Positive prices and coordinates.** Each transaction price $S_i(t)$ is strictly positive and initially $S_i(0) = F > 0$. A strictly increasing C^1 coordinate $q_i = \phi_i(S_i)$ preserves price priority. The analysis concerns deterministic admissible execution paths.

(M2) **Trade and execution accounting.** Executed share change is the signed quantity $\mathrm{d}\theta_i$; unsigned volume counts its absolute increments, not both counterparties twice. Executions obey marginal-price settlement $\mathrm{d}X_i = S_i\,\mathrm{d}\theta_i$. Matching uses price priority conditional on an order book; the effective price response is a separately imposed, positive, locally C^2 function $g(S)$, such that $\mathrm{d}S_i|_{\text{trade}} = g(S_i)\,\mathrm{d}\theta_i$.

(M3) **No fundamental news and resilience.** The benchmark F stays fixed. Between trades, prices recover toward F at a common constant rate $\rho > 0$: $\dot{S}_i = -\rho(S_i - F)$ in the absence of execution. The response function g and ρ are *model inputs*, not implications of price priority.

(M4) **Paired-dollar transaction network.** Each permitted transaction moves an equal executed-dollar amount between a source asset and a destination asset: $\mathrm{d}\boldsymbol{X} = B\,\mathrm{d}\boldsymbol{z}$, where B is the oriented incidence matrix of the complete bidirectional graph, $\mathrm{d}\boldsymbol{z} \geq 0$, and $\mathbf{1}^\top B = 0$. This describes a chosen paired-reallocation submarket; zero-sum aggressive order pressure is not a general fact about all real markets.

(M5) **Admissible controls and terminal closure.** Paths are piecewise continuous and of finite execution variation, with finitely many marginally settled instantaneous pulses allowed; every price remains positive and required sale inventory is available without short sales. The terminal comparison imposes $X_i(T) = 0$ for *each* asset, not $\theta_i(T) = \theta_i(0)$. Where resource caps are claimed, identical total paired-dollar turnover and per-asset starting share limits are imposed across markets.

Why the dollar ledgers close. The complete two-sided graph in (M4) represents an idealized portfolio-switching pool, or a cashless reallocation sleeve whose executed selling proceeds must immediately purchase other risky assets. It is not an identity for ordinary exchange-wide order flow: in a general equity market a cash account or additional counterparties may absorb proceeds. We impose the stronger terminal condition $X_i(T) = \int_0^T S_i \, \mathrm{d}\theta_i = 0$ to study residual price displacement after *every* asset's first-order net-dollar reallocation has canceled. This is a research normalization, not the usual definition of a trading round trip. A share-flat round trip instead imposes $\theta_i(T) = \theta_i(0)$, and the two constraints have different implications. In particular, from (2.11) below, dollar closure with terminal displacement or positive recovery dissipation requires consumption of initially available share inventory: $F[\theta_i(0) - \theta_i(T)] = E(S_i(T)) + \rho \int_0^T (S_i - F)^2/g(S_i) \, \mathrm{d}t$. Terminal dollar closure does not restore the number of shares. This distinction is essential to the barrier results.

Finite-variation paths and marginally settled pulses. For absolutely continuous controls the differential notation is classical. At a jump we do *not* use a left-endpoint or right-endpoint Stieltjes product $S_i \Delta\theta_i$. Instead physical time is held fixed and each directed edge $e : i \to j$ with dollar mass $z > 0$ is executed along a specified within-event parameter $a \in [0, z]$:

$$\frac{\mathrm{d}\widetilde{S}_i}{\mathrm{d}a} = -\frac{g(\widetilde{S}_i)}{\widetilde{S}_i}, \quad \frac{\mathrm{d}\widetilde{S}_j}{\mathrm{d}a} = \frac{g(\widetilde{S}_j)}{\widetilde{S}_j}, \quad \frac{\mathrm{d}\widetilde{\theta}_i}{\mathrm{d}a} = -\frac{1}{\widetilde{S}_i}, \quad \frac{\mathrm{d}\widetilde{\theta}_j}{\mathrm{d}a} = \frac{1}{\widetilde{S}_j}. \tag{2.1}$$

The resulting endpoint flows satisfy $\Delta X_i = -z$, $\Delta X_j = z$ and $\Delta\Phi(S_i) = -z$, $\Delta\Phi(S_j) = z$. Equivalently, a single-share impulse is evaluated by solving $\mathrm{d}\widetilde{S}/\,\mathrm{d}\xi = g(\widetilde{S})$ and integrating $\int \widetilde{S} \, \mathrm{d}\xi$ along that response trajectory. Jumps are thus *flow-completed* (Marcus-type) updates preserving the ordinary chain rule. Finite pulses on distinct edges are processed in their prescribed order. A pulse is admissible only while all intermediate prices remain strictly positive. These rules specify the interpretation of all jump expressions in what follows.

From the conditional matching map to an effective response. Fix a frozen book $\mathcal{B}$ with locally positive continuous share-depth density $\rho_S(S; \mathcal{B})$. A one-sided sweep whose signed marginal share size is $\mathrm{d}\theta$ obeys the local relation

$$\mathrm{d}\theta = \rho_S(S; \mathcal{B}) \, \mathrm{d}S, \qquad \rho_q(q; \mathcal{B}) \, \mathrm{d}q = \rho_S(S; \mathcal{B}) \, \mathrm{d}S. \tag{2.2}$$

The exact discrete version uses the generalized inverse of the book's cumulative depth. In the continuous approximation, changing $q = \phi(S)$ changes the density numerically through the Jacobian, not the matching priority. The relation refers to the marginal executed price. A dynamic midpoint, replenished book, and interasset cross-impact are *not* derived by coordinate transformation. Assumption (M2) instead chooses $g(S) = 1/\rho_S^{\mathrm{eff}}(S)$ as an *effective*, state-dependent response for the rest of the paper.

The resulting N-asset equation. Under the full market specification, the level-price law is

$$\boxed{\mathrm{d}S_i(t) = -\rho\big(S_i(t) - F\big) \, \mathrm{d}t + g(S_i(t)) \, \mathrm{d}\theta_i(t), \qquad \mathrm{d}X_i = S_i \, \mathrm{d}\theta_i.} \tag{2.3}$$

Choose the strictly monotone generalized coordinate $q_i = \log(S_i/F)$. Conversion to paired dollars gives

$$\mathrm{d}q_i\big|_{\mathrm{trade}} = \frac{g(S_i)}{S_i^2} \, \mathrm{d}X_i. \tag{2.4}$$

Combining this conversion, transaction pairing, and resilience yields the full multiasset system

$$\boxed{\mathrm{d}\boldsymbol{q} = -\rho(\mathbf{1} - e^{-\boldsymbol{q}})\,\mathrm{d}t + D(\boldsymbol{q})B\,\mathrm{d}\boldsymbol{z}, \qquad D_{ii}(\boldsymbol{q}) = \frac{g(Fe^{q_i})}{F^2 e^{2q_i}}.} \tag{2.5}$$

The exponential acts componentwise. This log coordinate is convenient for measuring losses, but executed dollars admit a simpler *response-adapted* coordinate. Let $y_i = \Phi(S_i)$ with Φ as in (1.1). It is strictly increasing and satisfies $\Phi(F) = 0$. Moreover, among differentiable coordinates ψ normalized by $\psi(F) = 0$, the requirement that each marginal trade obey $\mathrm{d}\psi(S)|_{\mathrm{trade}} = \mathrm{d}X$ forces $\psi'(S)g(S) = S$ and thus uniquely selects $\psi = \Phi$. Define the restoring response

$$\mathcal{R}(y) = \frac{S(S-F)}{g(S)}, \qquad S = \Phi^{-1}(y).$$

The entire multiasset market then has the *trade-normalized form*

$$\boxed{\mathrm{d}\boldsymbol{y} = B\,\mathrm{d}\boldsymbol{z} - \rho\mathcal{R}(\boldsymbol{y})\,\mathrm{d}t, \qquad \mathbf{1}^\top B = 0,} \tag{2.6}$$

where $\mathcal{R}$ acts componentwise and the execution rule applies also to flow-completed pulses. Trading is therefore an additive zero-sum input; all nonlinear price-state dependence has moved into the restoring term and the inverse map $S = \Phi^{-1}(y)$. The transformed state space is $\mathcal{J}^N$, where $\mathcal{J} = \Phi((0,\infty))$. For $g_p(S) = \eta(S/F)^p$, the range is $(-F^2/[\eta(2-p)], \infty)$ for $p < 2$, $\mathbb{R}$ for $p = 2$, and $(-\infty, F^2/[\eta(p-2)])$ for $p > 2$. Thus additive trading inputs do not remove strict price positivity or execution constraints. This change of coordinates does not alter which price states are reachable. With $\rho = 0$ the traded sum of generalized coordinates is exactly conserved.

Definition 2.2 (Terminal systemic drawdown)**.** Let $L_i(T) = \log[F/S_i(T)]$ and $\overline{L}(T) = N^{-1}\sum_i L_i(T)$. The kth-largest loss, where $k = \lceil \alpha N \rceil$, is $Q_\alpha(T)$. Define

$$\boxed{\mathcal{C}_\alpha(T) = \min\{\overline{L}(T), Q_\alpha(T)\}.} \tag{2.7}$$

For $\delta > 0$, a *terminal crisis* occurs when $\mathcal{C}_\alpha(T) \geq \delta$. Thus the mean loss and at least k individual losses must each reach δ. The target is selected independently of the price-response function.

The discrete coverage rule is material: for $N = 3$, $\alpha = 0.9$ means $k = 3$, so it requires a fall in all three assets. More generally, the contrast $k = N$ versus $k < N$ determines whether a compensating asset exists at all. This same distinction appears first in the A/B/C classification and later in the multiday cross-asset capacity inequality. A first-passage event $\tau_{\alpha,\delta} = \inf\{t : \mathcal{C}_\alpha(t) \geq \delta\}$ is different from a *terminal* crisis under dollar closure; we study the latter.

Definition 2.3 (Turnover, share inventory and terminal critical budget)**.** For an admissible path γ, set

$$M(\gamma) = \sum_e \int_0^T \mathrm{d}z_e, \qquad H_i^-(\gamma) = \int_0^T (-\,\mathrm{d}\theta_i)_+. \tag{2.8}$$

Under allowed initial inventory H_i and terminal dollar closure, define

$$M_*^{\mathrm{end}} = \inf\{M(\gamma) : X_i(T) = 0\ \forall i,\ \mathcal{C}_\alpha(T) \geq \delta,\ \gamma \text{ admissible}\}, \tag{2.9}$$

with $\inf \varnothing = +\infty$. A paired dollar transfer is counted once in M, although it appears with opposite signs in two asset ledgers.

The precise no-profitable-round-trip property. Define

$$E(S) = \int_F^S \frac{u - F}{g(u)}\,\mathrm{d}u \geq 0. \tag{2.10}$$

For any admissible path beginning at F the execution-cost identity is

$$\boxed{\int_0^T S_i\,\mathrm{d}\theta_i = F[\theta_i(T) - \theta_i(0)] + E(S_i(T)) + \rho\int_0^T \frac{(S_i - F)^2}{g(S_i)}\,\mathrm{d}t.} \tag{2.11}$$

Thus a *share-flat* round trip has nonnegative executed cost under (M2)–(M3). The paper's terminal condition $X_i(T) = 0$ instead closes executed *dollars*, not shares: a terminally dollar-flat path may consume initial inventory. Appendix A proves the accounting identity and its consequences for admissible controls.

Comparison with share-flat trading. The terminal dollar condition and the standard share-flat condition are not interchangeable. For example, with $g \equiv \eta$, a single asset purchased by a sufficiently small $a > 0$ shares at time 0 and sold by a shares at time T has $\theta(T) = \theta(0)$ but terminal price $S(T) = F - \eta a(1 - e^{-\rho T}) < F$; its net executed dollars are not zero and the positive cost in (2.11) is paid through the cash account. Conversely, under (M4), requiring every asset to be share-flat *as well as* terminally dollar-flat would force $\sum_i [E(S_i(T)) + \rho\int_0^T (S_i - F)^2/g(S_i)\,\mathrm{d}t] = 0$. Thus there can be no positive-time dissipative price excursion under both closures together. The dollar-flat protocol deliberately isolates an economic question distinct from standard share-round-trip execution.

3 Structural barriers to market crashes

The normal form (2.6) supplies a conserved scalar during trading. Write $H(\boldsymbol{S}) = \mathbf{1}^\top \boldsymbol{y} = \sum_i \Phi(S_i)$. Summing over assets cancels every matched dollar transfer and leaves only recovery:

$$\boxed{\mathrm{d}H = -\rho\sum_{i=1}^N \mathcal{R}(y_i)\,\mathrm{d}t = -\rho\sum_{i=1}^N \frac{S_i(S_i - F)}{g(S_i)}\,\mathrm{d}t.} \tag{3.1}$$

The scalar H is a trade invariant, not a conserved quantity under the full recovery dynamics. Whether the recovery drift can carry H through a crisis boundary depends on the curvature of Φ. Introduce the response elasticity

$$e(S) = \frac{Sg'(S)}{g(S)}. \tag{3.2}$$

The derivatives of Φ and $G(x) = \Phi(Fe^x)$ are

$$\Phi''(S) = \frac{1 - e(S)}{g(S)}, \qquad G''(x) = \frac{S^2}{g(S)}\,[2 - e(S)], \quad S = Fe^x. \tag{3.3}$$

The distinction between these two curvatures tracks the definition of systemic loss. A barrier to *every asset* being below its benchmark is weaker than a barrier to our functional, which also involves the *average* log loss. Thus the same coordinate can produce two different exclusion statements.

Theorem 3.1 (Barrier to a decline of every asset)**.** *If $e(S) \geq 1$ at all prices attainable by admissible paths, then $S_i(t) < F$ for every i is impossible at any time t. In particular, no positive crisis with $k = N$ is reachable.*

Proof. Equation (3.3) makes Φ concave. Its tangent inequality at S is $\Phi(F) \le \Phi(S) + \Phi'(S)(F-S)$, hence $(S-F)\Phi'(S) \le \Phi(S)$. Consequently (3.1) implies $\dot{H} \ge -\rho H$ between trades; pulses preserve H. Since $H(0) = 0$, comparison gives $H(t) \ge 0$. If every $S_i(t) < F$, strict monotonicity and $\Phi(F) = 0$ instead imply $H(t) < 0$, a contradiction. □

Theorem 3.2 (Barrier to a systemic crisis of any breadth)**.** *If $e(S) \ge 2$ on the admissible price range, then for every admissible control, $N \ge 2$, $\alpha \in (0,1]$ and $t \in [0,T]$,*

$$\boxed{\mathcal{C}_\alpha(t) \le 0.} \tag{3.4}$$

Thus no strictly positive systemic terminal-crisis threshold can be reached.

Proof. The preceding theorem gives $H(t) \ge 0$. The second identity in (3.3) makes $G(x) = \Phi(Fe^x)$ concave and strictly increasing. Jensen's inequality yields

$$0 \le \frac{H(t)}{N} = \frac{1}{N}\sum_i G(q_i(t)) \le G\left(\frac{1}{N}\sum_i q_i(t)\right).$$

As $G(0) = 0$ and G is strictly increasing, $\sum_i q_i(t) \ge 0$. Hence the mean loss $\overline{L}(t) = -N^{-1}\sum_i q_i(t)$ is nonpositive; Definition 2.2 implies (3.4). □

The two theorems locate states that *no* admissible trade program can reach under their respective assumptions. They are global *sufficient* conditions, not converses and not necessary conditions for stability. If the inequalities hold only near F, the exclusions apply only while the entire path stays there; limited inventory, trading links, or time may create further barriers. The next question is consequently constructive: *does relaxing the structural condition actually change the reachable set?* The A/B/C comparison answers this within one response family, rather than interpreting the failure of a sufficient condition as proof of a crash.

4 Three market regimes under common constraints

The curvature conditions above establish what is forbidden; they do not tell us what becomes possible when an inequality fails. To exhibit the distinction under matched assumptions rather than unrelated models, we change only the response shape in the family

$$g_p(S) = \eta\left(\frac{S}{F}\right)^p, \qquad p \in \{2,1,0\}, \quad \eta > 0. \tag{4.1}$$

All three markets have identical N, T, F, ρ, η, identical initial prices and initial *impact*, the same paired-dollar trading network, and, when compared numerically, the same turnover and inventory caps. They differ only in the exponent p. In the log-dollar equation (2.5),

$$D_{ii}^{(p)}(q_i) = \frac{\eta}{F^2}e^{(p-2)q_i}, \qquad e(S) = p. \tag{4.2}$$

The same coordinate formula integrates explicitly for every exponent:

$$\Phi_p(S) = \begin{cases} \dfrac{F^p}{\eta(2-p)}(S^{2-p} - F^{2-p}), & p \ne 2, \\ \dfrac{F^2}{\eta}\log(S/F), & p = 2. \end{cases}$$

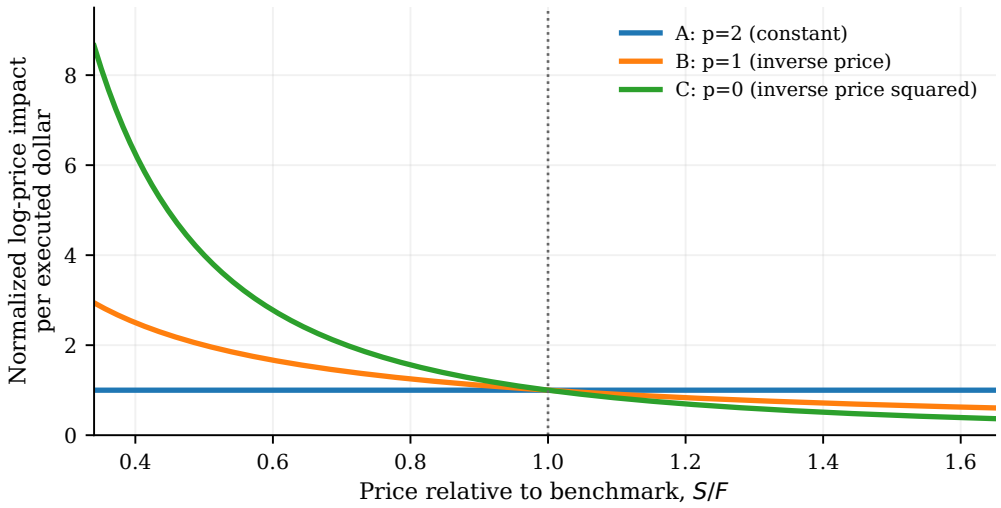


Figure 1: A single response family, three price sensitivities. The plot shows the model-implied instantaneous *log-price change per executed dollar*, normalized by its common initial value. All curves coincide at $S = F$. It is a theoretical comparison, not a fit to observed order books.

Consequently, paired trades conserve a quantity proportional to $\sum_i S_i^2$, $\sum_i S_i$, or $\sum_i \log S_i$ for $p = 0, 1, 2$, respectively. Recovery evolves these quantities between trades. Figure 1 shows the corresponding instantaneous response per executed dollar: below F, the $p = 0$ response amplifies dollar trades most strongly.

Theorem 4.1 (A/B/C reachability regimes)**.** *Let $N \geq 2$ and $k = \lceil \alpha N \rceil < N$. Compare the three markets in (4.1) under the same model protocol.*

(A) ***Barrier market*** *($p = 2$):* $\mathcal{C}_\alpha(t) \leq 0$ *for every admissible path and every horizon.*

(B) ***Inventory-bounded market*** *($p = 1$): positive but sufficiently small terminal systemic drawdowns are reachable. If every asset starts with a fixed finite inventory cap H_0 and short selling is forbidden, the terminal crisis severity is bounded above independently of the trade sequence.*

(C) ***Severely reachable market*** *($p = 0$): for every fixed $N \geq 2$ and $k < N$, there exists a sufficiently long but finite horizon T and finite constants M_0, H_0 such that, for* every *finite $\delta > 0$, an admissible positive-price path with $M \leq M_0$, $H_i^- \leq H_0$, and $X_i(T) = 0$ attains $\mathcal{C}_\alpha(T) > \delta$.*

The constants in (C) may depend on N, T, F, ρ, η, but not on δ. They do not remain uniform as $N \to \infty$.

Proof. (A) follows from Theorem 3.2. (B) follows from the two-pulse construction in Proposition 4.3 and the inventory inequality (4.3) below. (C) is established constructively in Appendix C: $N - 1$ targets are driven arbitrarily close to zero, while one asset serves as a cash-flow compensator. The control uses positive prices, fully paired marginal-dollar executions, a partition of recovery time, and terminal closure of every X_i. Its explicit turnover and inventory bounds are independent of the target final price $\varepsilon \downarrow 0$. □

Under the same executed-dollar and recovery protocol, the three response functions produce distinct attainable-state geometries. Market A supplies a global invariant-region barrier; Market B shows that nonzero partial-coverage losses need not permit arbitrarily large ones when stock inventory is finite; and Market C supplies a controlled severe-loss witness. The condition $k < N$ in (C) is essential: at least one asset must be available to carry the compensating flow. This will become an explicit *capacity* constraint, rather than merely an existence requirement, once finite gross share rates are imposed.

For $p = 1$, the inventory potential from (2.10) is

$$E_1(Fe^{-L}) = \frac{F^2}{\eta}\left(L + e^{-L} - 1\right). \tag{4.3}$$

By (2.11), terminal dollar closure and no short selling imply $E_1(S_i(T)) \le FH_0$. Since $L+e^{-L}-1 \to \infty$ as $L \to \infty$, no fixed finite inventory can support arbitrarily severe losses of the requisite k assets. By contrast, for $p = 0$, $E_0(S) = (S - F)^2/(2\eta)$ remains bounded as $S \downarrow 0$; this makes the construction in (C) possible but does not by itself prove reachability.

Proposition 4.2 (Local all-asset decline)**.** *If g is C^3 near F and $e(F) < 1$, then for every $N \ge 2$ there exists a terminally dollar-closed, positive-price, small-amplitude three-pulse control under which* every *$S_i(T) < F$. With $c = g(F)/F^2$ and $r = e^{-\rho T/2}$, the losses of a zero-sum impulse direction $\boldsymbol{u}$ with every $u_i \neq 0$ satisfy*

$$L_i(T) = \frac{1 - e(F)}{2} c^2 h^2 u_i^2 r(1 - r^2) + O(h^3) > 0 \tag{4.4}$$

for sufficiently small $h > 0$.

Proof. Apply the executed-dollar pulses $h\boldsymbol{u}$, $-(1 + r)h\boldsymbol{u}$, $rh\boldsymbol{u}$ at $0, T/2, T$. They are pairwise dollar-balanced and sum to zero for every asset. Taylor expansion in the coordinate $y = \Phi(S)$ gives (4.4); details appear in Appendix B. □

Proposition 4.3 (Local partial-coverage decline)**.** *Let $k < N$, $e_0 = e(F) < 2$, and $r = e^{-\rho T} < 2 - e_0$. Select a zero-sum dollar direction with at least k positive components and execute $h\boldsymbol{u}$ at time 0 and $-h\boldsymbol{u}$ at time T. Then $\mathcal{C}_\alpha(T) > 0$ for all sufficiently small $h > 0$, and*

$$\overline{L}(T) = \frac{c^2 h^2}{2N}(1 - r)(2 - e_0 - r)\sum_i u_i^2 + O(h^3). \tag{4.5}$$

Proof. The individual losses satisfy $L_i(T) = c(1-r)hu_i + O(h^2)$, while the first-order term vanishes when averaged because $\sum_i u_i = 0$. The strictly positive quadratic coefficient in (4.5) proves the result. See Appendix B. □

The next two propositions explain the local end of the classification. A two-pulse path can generate a partial-coverage loss under a less restrictive elasticity threshold; a three-pulse path with an additional cancellation step can generate a small decline of *every* asset when $e(F) < 1$. Both are local constructions, not substitutes for the global severe-loss argument in Market C.

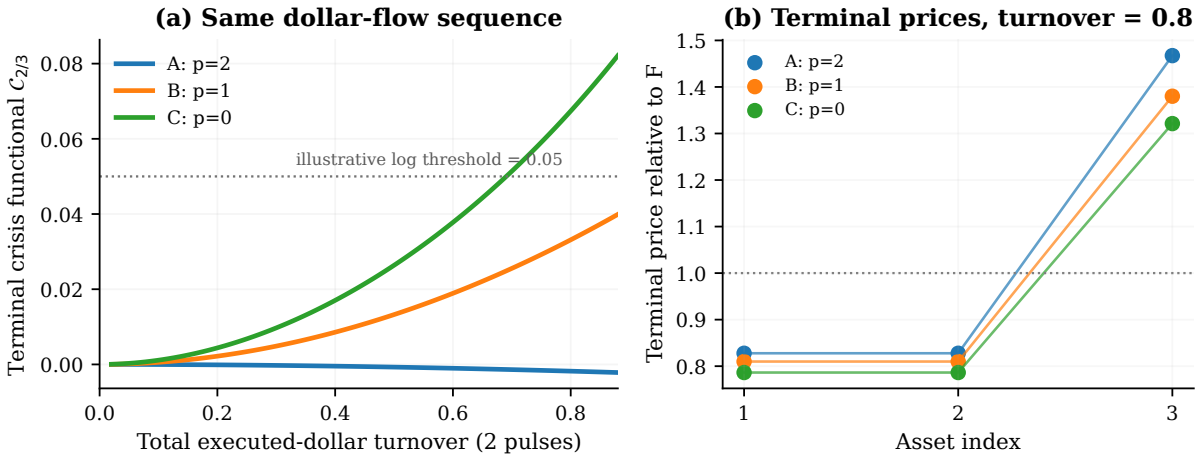


Figure 2: An identical trading experiment across A/B/C response functions. Two terminally offsetting executed-dollar pulses act on $N = 3$ assets with $k = 2$, $F = \eta = \rho = 1$, and $T = 3$. Left: the measured *terminal* crisis functional as paired-dollar turnover increases; right: terminal prices under the same total turnover $M = 0.8$. This local illustration is distinct from the severe-crash existence construction. The dotted threshold is illustrative, not calibrated.

A concrete common-budget comparison. For $N = 3$, $k = 2$, $F = \eta = \rho = 1$, $T = 3$, common turnover cap $M_0 = 3$, and per-asset initial inventory cap $H_0 = 8.2$, the $p = 2$ market cannot reach any positive crisis; the $p = 1$ market can reach a small positive crisis but not the threshold $\delta = 10$ (it would require over 9 shares of selling by any qualifying asset); a constructive $p = 0$ path reaches $\mathcal{C}_{2/3}(T) \approx 13.32$ with turnover approximately 2.62 and a maximum per-asset cumulative share sale below 3.94. The large δ is chosen as a rigorous stress-test separating bounded from unbounded *log* loss; it is not a representative empirical crash size. Validation code accompanies the source files.

4.1 A sharp boundary when every asset must fall

Market C relies on $k < N$. To see why the breadth condition matters beyond a technical assumption, keep *the same constant-impact response* but change the target to $k = N$. There is then no nonqualifying asset available to compensate the other $N - 1$ terminal losses. Rather than an arbitrary-severity construction, the model now yields an exact horizon-dependent ceiling, even before a fixed execution budget is imposed.

Theorem 4.4 (Sharp all-asset horizon threshold)**.** *Under Assumption 2.1 with $g(S) = \eta$ and $k = N$, define*

$$\boxed{\delta_{\mathrm{crit}}(N, T) = \log \frac{\sqrt{N}}{1 + (\sqrt{N} - 1)e^{-\rho T}}.} \tag{4.6}$$

For every $\delta \in (0, \delta_{\mathrm{crit}}(N, T))$, some admissible terminally dollar-closed control reaches $\mathcal{C}_1(T) \geq \delta$ with finite turnover and sufficiently large finite inventory. If $\delta \geq \delta_{\mathrm{crit}}(N, T)$, no such positive-price path exists. Thus δ_{crit} is the unattained supremum *of the all-asset common log drawdown.*

Proof sketch. Set $R(t)^2 = \sum_i S_i(t)^2$. Under constant response, any paired-dollar trade satisfies $\mathrm{d}R^2|_{\mathrm{trade}} = 2\eta \sum_i \mathrm{d}X_i = 0$. Between trades,

$$\dot{R} = -\rho \left(R - F \frac{\sum_i S_i}{R} \right).$$

For positive prices, $\sum_i S_i/R > 1$, whence $R(T) > F\{1 + (\sqrt{N} - 1)e^{-\rho T}\}$. If each asset ends below $Fe^{-\delta}$, then $R(T) \le F\sqrt{N}e^{-\delta}$, proving necessity. For sufficiency, rotate the price vector through N positive spherical directions with one high coordinate, allotting each asset equal recovery weight. Each reorientation is a zero-sum paired-dollar pulse, and recovery-stage paired controls keep the direction fixed. Equal terminal prices and equal recovery integrals, together with $\sum_i X_i(T) = 0$, give $X_i(T) = 0$ for every i. Moving the positive low coordinates toward zero approaches the bound continuously. A full constructive proof is in Appendix D. □

5 How much trading is necessary?

The classification so far answers *whether* the market admits a loss state. A feasible market need not reach it with the resources available to its participants, however. We therefore shift from price-state geometry to a second filter: *how much executed dollar turnover and how many shares must be sold to attain a specified breadth and severity?* The following bounds work for general positive responses and complement, rather than replace, the A/B/C reachability results.

Theorem 5.1 (Execution-resource lower bounds)**.** *Suppose an admissible terminal path satisfies $\mathcal{C}_\alpha(T) \ge \delta > 0$. Write $k = \lceil \alpha N \rceil$ and $s = Fe^{-\delta}$. Besides Φ, define the monotone share-response coordinate $\Psi(S) = \int_F^S \mathrm{d}u/g(u)$. Then the necessary dollar-turnover and total share-sale bounds are simply the respective coordinate distances to s:*

$$\boxed{M \ge -k\Phi(s) = k\int_s^F \frac{u}{g(u)}\,\mathrm{d}u, \qquad \sum_i H_i^- \ge -k\Psi(s) = k\int_s^F \frac{\mathrm{d}u}{g(u)}.} \tag{5.1}$$

Moreover, if $X_i(T) = 0$, no short sales are allowed, and the initial inventory of any qualifying asset does not exceed H_i, then

$$\boxed{\int_s^F \frac{F-u}{g(u)}\,\mathrm{d}u \le FH_i.} \tag{5.2}$$

These inequalities are necessary, not sufficient, and need not be sharp.

Proof. For a terminally qualifying asset, the final price is at most $s < F$. Both transformed coordinates increase with price, and recovery acts toward their zero values at F. Consequently the cumulative negative executed-dollar and share variations must be at least $-\Phi(s)$ and $-\Psi(s)$, respectively. Summing over k distinct qualifying assets proves (5.1) because each paired execution is counted exactly once as a sale across assets. Equation (5.2) follows from the nonnegative right-hand side of (2.11) with $X_i(T) = 0$ and $\theta_i(0) - \theta_i(T) \le H_i$. The detailed variation argument is given in Appendix A. □

The resource bounds are geometric distances in two monotone coordinates, not optimal trade schedules. Trading direction, recovery time, and the buying capacity of receiving assets impose further requirements; we turn to these execution restrictions next.

If N is increased by *replicating identical market depths* and α, δ are held fixed, then $k = \Theta(N)$ and (5.1) entails $M_*^{\mathrm{end}} = \Omega(N)$ whenever the target is reachable. This is not a universal capitalization-scaled law for actual exchanges: increasing N under this experiment also increases aggregate depth.

The low-price integrals discriminate different kinds of resource scarcity:

$$\int_0^F \frac{u}{g(u)}\,\mathrm{d}u, \qquad \int_0^F \frac{\mathrm{d}u}{g(u)}. \tag{5.3}$$

For a local power law $g(u) \asymp u^p$, the dollar integral converges only if $p < 2$, and the share integral only if $p < 1$. These integral criteria separate resource feasibility from the structural and dynamic conditions for reachability.

6 Finite trading rates and daily price limits

A structurally reachable loss may require trades that cannot be completed at finite speed. We therefore restrict the control set by gross executed-share rates and single-session price bands. These restrictions preserve the barrier theorems, while positive reachability requires a new continuous construction rather than the pulse controls used in Market C.

The imposition of finite execution speed builds on rate-constrained trading models such as Yang [18], while our goal remains the reachability of a collective loss threshold rather than the optimum of an agent's execution objective.

6.1 Two separate execution restrictions

For absolutely continuous nonnegative paired-dollar edge rates $\dot{z}_e(t)$, define an asset's *gross executed share rate* by

$$r_i^{\mathrm{gross}}(t) = \frac{1}{S_i(t)} \sum_e |B_{ie}| \dot{z}_e(t), \qquad 0 \le r_i^{\mathrm{gross}}(t) \le \bar{v}, \qquad i = 1, \dots, N. \tag{6.1}$$

This bound counts buy and sell executions *before* netting. In particular $|\dot{\theta}_i| \le r_i^{\mathrm{gross}} \le \bar{v}$ and no Dirac-sized trade is admissible. Bounding only the signed net share rate would allow simultaneous offsetting gross transactions of arbitrary size and would not constrain executed turnover.

Independently, a stylized *one-session price-limit rule* restricts every executed price to

$$F(1 - \beta_-) \le S_i(t) \le F(1 + \beta_+), \qquad 0 < \beta_- < 1, \quad \beta_+ > 0. \tag{6.2}$$

The control must keep its price path inside this interval; an attempted trade that would cross the limit is *inadmissible*, not reflected mechanically back into the band. Equation (6.2) is an abstraction: TSE price limits are predominantly absolute yen intervals based on the prior close or base price, vary with the price level, and may be widened across sessions [12]. We do not model queued unfilled orders, special quotes, limit-hit auction allocation, or next-day reopening.

Proposition 6.1 (Rate and price-band exclusion certificates)**.** *Suppose the common impact function is positive and $g(S) \le G$ on the whole admissible price region. Under the cap (6.1), every asset obeys*

$$S_i(t) \ge F - \frac{G\bar{v}}{\rho}(1 - e^{-\rho t}), \qquad 0 \le t \le T. \tag{6.3}$$

Under the additional price band (6.2), put

$$S_{\mathrm{floor}} = \max\left\{ F(1 - \beta_-),\ F - \frac{G\bar{v}}{\rho}(1 - e^{-\rho T}) \right\} > 0. \tag{6.4}$$

Then for any breadth α and any admissible control,

$$\boxed{\mathcal{C}_\alpha(t) \le \log \frac{F}{S_{\text{floor}}} \quad \text{for all } t \le T.} \tag{6.5}$$

In particular, a necessary finite-speed condition to reach threshold $\delta > 0$ is

$$\boxed{\bar{v} \ge \frac{\rho F(1-e^{-\delta})}{G(1-e^{-\rho T})}.} \tag{6.6}$$

A one-session lower price limit independently excludes any $\delta > -\log(1-\beta_-)$. These are necessary *constraints, generally not sharp reachability thresholds. For comparison, if a* hypothetical *multi-session regime resets a proportional lower band to each preceding close, then over D sessions the elementary bound is only $S_i(T_D) \ge F\prod_{d=1}^{D}(1-\beta_{-,d})$ and $\mathcal{C}_\alpha(T_D) \le -\sum_{d=1}^{D}\log(1-\beta_{-,d})$. Thus a daily band cannot be interpreted as a permanent, fixed-baseline ban on a cumulative crash. Actual TSE limits are price-tiered yen rules, not these proportional resets.*

Proof. The variation-of-constants formula for $\dot{S}_i = -\rho(S_i - F) + g(S_i)\dot{\theta}_i$ and $|\dot{\theta}_i| \le \bar{v}$ yields $S_i(t) \ge F - G\bar{v}\int_0^t e^{-\rho(t-s)}\,\mathrm{d}s$, proving (6.3). The price band supplies the second floor; every individual log loss is at most $\log(F/S_{\text{floor}})$, and so are its cross-sectional mean, upper-breadth threshold, and their minimum. Rearranging the rate-only floor is necessary for an individual loss to reach δ. □

The per-asset comparison controls how far one price can move. But our crisis functional needs many prices to move together, and the transaction protocol counts *paired* dollars. The turnover and terminal-dollar-closure protocol therefore provide two further certificates that translate the same loss threshold into market-wide execution requirements. Since each paired dollar is counted twice in gross share execution,

$$M = \frac{1}{2}\int_0^T \sum_i S_i(t) r_i^{\text{gross}}(t)\,\mathrm{d}t \le \frac{NF(1+\beta_+)\bar{v}T}{2}. \tag{6.7}$$

Together with Theorem 5.1, this gives the necessary joint constraint

$$\boxed{\bar{v} \ge \frac{2k}{NF(1+\beta_+)T}\int_{Fe^{-\delta}}^{F}\frac{u}{g(u)}\,\mathrm{d}u.} \tag{6.8}$$

For a closed executed-dollar ledger, the nonnegative cost identity (2.11) also yields

$$E(S_i(T)) \le F\bar{v}T, \qquad \bar{v} \ge \frac{E(Fe^{-\delta})}{FT} \quad \text{for every qualifying asset.} \tag{6.9}$$

Bounds (6.6), (6.8), and (6.9) are distinct necessary tests; satisfying them does not construct a crisis path.

A naive per-asset rate restriction would appear to produce a first-order $O(\bar{v})$ limit on systemic losses. Dollar pairing is stronger: at small rates, the leading-order movements of different assets cancel in the *market mean*. This is why the next constant-impact result is quadratic rather than merely linear, and why it is naturally compared with a small, continuous construction.

Theorem 6.2 (A quadratic small-rate crisis barrier for constant impact)**.** *Let $g(S) \equiv \eta > 0$, assume dollar pairing $\sum_i S_i \dot{\theta}_i = 0$, and impose the gross share-rate cap (6.1). Write $d_T = 1 - e^{-\rho T}$ and*

$$b_T = \frac{\eta \bar{v}}{F\rho}\, d_T < 1. \tag{6.10}$$

Then, with or without terminal dollar closure and for every *$N \geq 2$ and breadth α,*

$$\boxed{\mathcal{C}_\alpha(T) \leq \left(\frac{\eta \bar{v}}{F\rho}\right)^2 \left[d_T - \rho T e^{-\rho T} + \frac{d_T^2}{2(1-b_T)^2}\right].} \tag{6.11}$$

In particular, the maximum possible positive crisis severity is $O(\bar{v}^2)$ as $\bar{v} \downarrow 0$. For fixed N, T and any positive price-band width, this order is attained *by the continuous control of Theorem 6.3, and hence its minimum required gross execution-rate cap obeys $\bar{v}_*(\delta) = \Theta(\sqrt{\delta})$ as $\delta \downarrow 0$ in the constant-impact market. The implied coefficients depend on the prescribed control and are not an exact optimization solution.*

Proof. Set $x_i(t) = S_i(t) - F$ and $m(t) = \sum_i x_i(t)$. Variation of constants gives $|x_i(t)| \leq \eta\bar{v}(1 - e^{-\rho t})/\rho$. Since $\sum_i S_i \dot{\theta}_i = 0$, $\sum_i \dot{\theta}_i = -F^{-1} \sum_i x_i \dot{\theta}_i$ and therefore $\dot{m} = -\rho m - (\eta/F) \sum_i x_i \dot{\theta}_i$. Integrating and bounding the gross rates yields

$$\frac{|m(T)|}{NF} \leq \left(\frac{\eta \bar{v}}{F\rho}\right)^2 [d_T - \rho T e^{-\rho T}]. \tag{6.12}$$

For $|x_i(T)|/F \leq b_T < 1$, Taylor's inequality $-\log(1+y) \leq -y + y^2/[2(1-b_T)^2]$ implies $\overline{L}(T) \leq |m(T)|/(NF) + b_T^2/[2(1-b_T)^2]$. Use $\mathcal{C}_\alpha \leq \overline{L}$ to obtain (6.11). The construction in Theorem 6.3 has $\mathcal{C}_1(T) = c_N a^2 + O(a^3)$ with $c_N > 0$, and its gross rate is $O(a)$; setting $a \propto \bar{v}$ gives a matching positive $\Omega(\bar{v}^2)$ witness while keeping a fixed positive band. This proves the stated small-rate exponent. □

Theorem 6.3 (Positive all-asset decline with bounded continuous execution)**.** *Assume $g \in C^3$ near F, $g(F) > 0$, $\rho > 0$, and $e_0 = Fg'(F)/g(F) < 1$. Fix* any *$N \geq 2$, horizon $T > 0$, rate cap $\bar{v} > 0$, one-session band widths $\beta_-, \beta_+ > 0$, and any fixed positive amount of initially available shares per asset. Then there exists an* absolutely continuous*, paired-dollar, per-asset gross-rate-bounded control, remaining strictly inside the band, with $X_i(T) = 0$ for each asset and*

$$\boxed{S_i(T) < F \text{ for every } i, \qquad \mathcal{C}_1(T) > 0.} \tag{6.13}$$

The attained loss can be arbitrarily small. This does not *assert that a prescribed fixed large threshold remains reachable under an arbitrarily tight cap or band.*

Construction and sign calculation. Choose a zero-sum vector $u \in \mathbb{R}^N$ with every $u_i \neq 0$. Put $\Delta = T/3$, $r = e^{-\rho\Delta}$ and set the bounded, three-bin signal

$$f(t) = \begin{cases} 1, & 0 \leq t < \Delta, \\ -(1+r), & \Delta \leq t < 2\Delta, \\ r, & 2\Delta \leq t \leq T, \end{cases} \qquad \dot{X}_i = a u_i f(t). \tag{6.14}$$

The dollar rates sum to zero, and $\int_0^T f(t)\,\mathrm{d}t = 0$. Direct integration also gives $\int_0^T e^{-\rho(T-t)} f(t)\,\mathrm{d}t = 0$. Write $h(S) = g(S)/S$ and expand $S_i(t) = F + a u_i y(t) + a^2 u_i^2 z(t) + O(a^3)$, where $\dot{y} = -\rho y + h(F) f$,

$y(0) = 0$, and $\dot{z} = -\rho z + h'(F)fy$, $z(0) = 0$. The weighted cancellation implies $y(T) = 0$. Since $h'(F)/h(F) = (e_0 - 1)/F$, integration by parts yields

$$\boxed{\log \frac{F}{S_i(T)} = \frac{\rho(1 - e_0)}{2F^2}\, a^2 u_i^2 \int_0^T e^{-\rho(T-t)} y(t)^2 \, \mathrm{d}t + O(a^3).} \tag{6.15}$$

The integral is strictly positive. Thus all prices fall for every sufficiently small $a > 0$. Represent each zero-sum dollar-rate vector by nonnegative source-to-destination edge rates with minimum gross flow, so $r_i^{\text{gross}} = |au_i f|/S_i$. The entire price path converges uniformly to F as $a \downarrow 0$, making both the finite rate cap, the price band, and any positive inventory availability valid for sufficiently small a. □

The construction complements the exclusion bound: a finite rate cap does not force all terminal prices back to the benchmark under dollar closure. It preserves local common-loss reachability while restricting the larger pulse-driven control set.

What survives when the restrictions are imposed separately? The global barrier for $e(S) \geq 2$ and the full-coverage exclusion for $e(S) \geq 1$ remain valid after restricting the set of admissible controls. For $p = 1$, a bounded two-stage buy-then-sell dollar-rate control still makes $k < N$ assets end below F: its per-asset signed dollars integrate to zero, while $\sum_i S_i(T) = NF$ and strict concavity of log makes the average log loss positive. In contrast, the *arbitrarily severe* construction for $p = 0$ in Theorem 4.1(C) is not a valid witness under (6.1), because it uses instantaneous sweeps. Under a fixed price band it is *mathematically impossible* for any one-session terminal loss to exceed $-\log(1-\beta_-)$. Theorem 6.3 proves what genuinely survives both restrictions: strictly positive, but not arbitrarily severe, local all-asset declines.

Table 1: How distinct execution restrictions change the $p = 0$ construction for a fixed T. "Local" means a sufficiently small positive terminal all-asset loss, not an arbitrary prescribed large loss.

Admissible controls	Proven survivability	What can no longer be inferred
Instantaneous sweeps, no band	Arbitrarily severe partial-coverage log loss, $k < N$	Does not model a bounded execution rate
Gross rate $\bar{v}$ only	Local all-asset loss for any $\bar{v} > 0$	Unbounded severe-loss construction not certified
Hard session band only	Local all-asset loss when band has nonzero width	$\delta > -\log(1 - \beta_-)$ impossible
Both restrictions	Local all-asset loss for positive caps/widths; rate and price-floor certificates	No general exact severe threshold or optimal execution rate

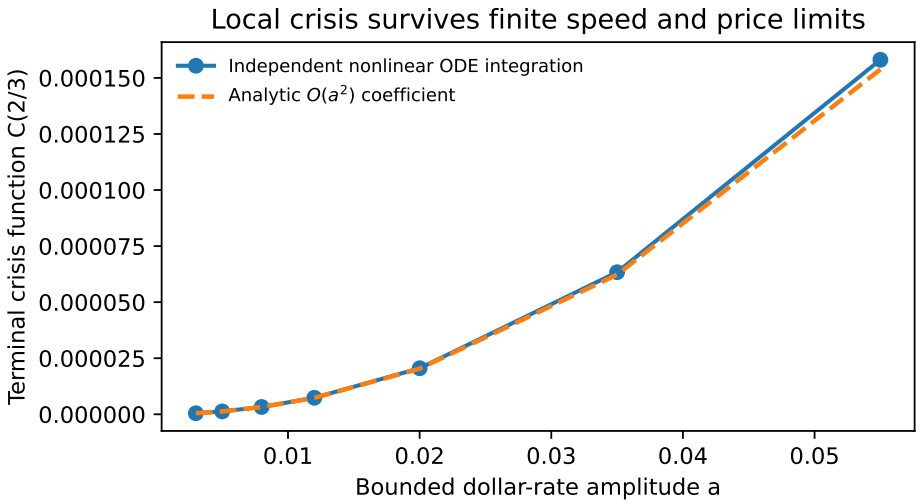


Figure 3: Quadratic local all-asset loss under a finite-rate control. Independently integrated nonlinear price paths for constant g converge to the analytical a^2 coefficient in (6.15). The plotted amplitude is sufficiently small to preserve positive prices; the admissibility checks for the highlighted $a = 0.035$ case are reported in the code bundle.

7 Consecutive limit-down closes across sessions

A one-session price floor limits *today's* loss, but the next session may begin from today's lower closing price. To extend the preceding reachability analysis across days, we must now keep track of *which* assets repeatedly reach their lower limits, how the crisis functional aggregates the resulting losses, and whether there is enough executed-dollar capacity on the other side of each sale. These are linked constraints, not separate illustrations: the target must clear a breadth threshold, every target must be sold fast enough, and the complement must accommodate the paired flow within its own rate and band.

The term *lower-limit close* here means the last executed price equals that session's floor; it is not a *locked limit-down order book* with unfilled sell queues. Queue dynamics are absent from our executed-trade state. The distinction matters when interpreting price-limit evidence such as Kim and Rhee [13], Chen et al. [5]. We proceed from a geometric rolling-band ceiling to a conditional crisis test, then to rate, conservation, and cross-asset capacity conditions, and finally to a certified feasible path.

7.1 Rolling bands and the exact cumulative accounting

Partition trading time into D sessions $[T_{d-1}, T_d]$, each of duration $\tau_d > 0$. Overnight calendar time does not contribute to the model clock; prices carry over continuously from the preceding close. In the stylized *proportional-reset* version, set $R_{i,d} = S_i(T_{d-1})$ and require

$$(1 - \beta_{i,d,-})R_{i,d} \le S_i(t) \le (1 + \beta_{i,d,+})R_{i,d}, \qquad t \in [T_{d-1}, T_d], \tag{7.1}$$

where $0 < \beta_{i,d,-} < 1$. A lower-limit close is $S_i(T_d) = (1 - \beta_{i,d,-})R_{i,d}$; a D-day streak means this equality holds for the *same* assets in D consecutive sessions. A streak in $k = \lceil \alpha N \rceil$ assets is not yet a systemic crisis under Definition 2.2: the market-average log loss must also reach the chosen threshold.

For a common constant lower fraction $q = 1 - \beta_-$, each streaking asset has $S_i(T_D) = Fq^D$ and loss $L_i(T_D) = D\lambda$ with $\lambda = -\log q$. Conversely, even without a streak, the state constraint ensures

$$\mathcal{C}_\alpha(T_D) \le D\lambda \tag{7.2}$$

for a common band and initial equal benchmark price. For heterogeneous bands the right side can be replaced by $\max_i \sum_{d=1}^{D}[-\log(1-\beta_{i,d,-})]$. These are *geometric ceilings*, not sufficient conditions for a dollar-paired trajectory to hit the floor.

In actual Tokyo Stock Exchange rules the daily absolute yen limits depend on the base-price tier, rather than a common percentage, and the base price can incorporate a special quote. The corresponding conceptual recursion is $R_{i,d+1} = R_{i,d} - \Delta_-(R_{i,d})$ for a daily lower-limit close, provided that the base reference equals the preceding close. Special procedures can widen the daily band after two specified consecutive limit-bound sessions with zero or minimal execution [12]. Our proportional-reset model deliberately omits this conditional market rule.

7.2 Cross-sectional crisis threshold under a closing-limit streak

The first gate is *definitional*. Consecutive lower-limit closes by k stocks satisfy the breadth requirement, but do not automatically satisfy the mean-loss requirement: the remaining assets may appreciate as they receive matched dollars. This is the same $k < N$ compensation mechanism that permits Market C's partial-coverage outcomes, now seen from the opposite direction. The following exact endpoint identity shows when the falling majority becomes a crisis under the *unchanged* functional, without yet constructing the trades.

Proposition 7.1 (A conditional systemic-crisis trigger test)**.** *Suppose the same $k = \lceil \alpha N \rceil < N$ assets close at Fq^D after D identical proportional-reset sessions, with $q \in (0,1)$. Let the remaining $N-k$ assets close at $Y_j > 0$, and write $\lambda = -\log q$. For any $0 < \delta \le D\lambda$, the* terminal *crisis threshold $\mathcal{C}_\alpha(T_D) \ge \delta$ holds if and only if*

$$\boxed{\prod_{j=1}^{N-k} \frac{Y_j}{F} \le \exp\big(kD\lambda - N\delta\big).} \tag{7.3}$$

If all other assets have individual log losses at most $D\lambda$ (or if $k > N/2$, so the $N-k$ other assets cannot fill the top k ranks), then $Q_\alpha(T_D) = D\lambda$ and the terminal crisis severity is exactly

$$\boxed{\mathcal{C}_\alpha(T_D) = \min\left\{ D\lambda, \frac{kD\lambda - \sum_{j=1}^{N-k} \log(Y_j/F)}{N} \right\}.} \tag{7.4}$$

In particular, with a single compensator ($k = N-1$), its terminal price must obey $Y/F \le \exp(kD\lambda - N\delta)$ if a crisis of threshold δ is to occur. These are exact conditional endpoint identities*, not sufficient conditions for a trade path realizing the stipulated closing prices.*

Proof. The k streaking assets have loss $D\lambda$, hence $Q_\alpha \ge D\lambda$. For $\delta \le D\lambda$, the coverage part of the crisis criterion is therefore automatic, and the remaining requirement is $A \ge \delta$. Substituting the terminal prices into $A = N^{-1}\sum_i \log(F/S_i)$ gives (7.3). If every other loss is at most $D\lambda$, or if $N - k < k$, the kth-largest loss is exactly $D\lambda$, proving (7.4). □

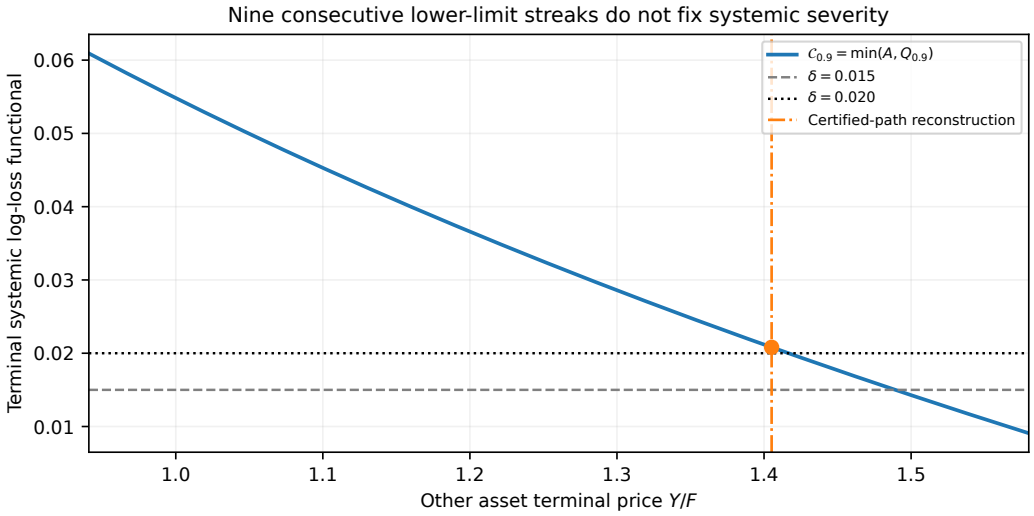


Figure 4: A closing-limit streak is not, alone, a systemic crisis. Conditional terminal severity for $N = 10$, nine assets closing twice at a 3% daily lower limit, versus the remaining asset's terminal price. The dotted marker is the constructed dollar-closed feasible path; 1.5% and 2.0% are log-loss thresholds, not daily price limits. The curve is an algebraic endpoint criterion, not a claim that every compensator price is dynamically reachable.

For the interval-certified $N = 10$, $k = 9$, $q = 0.97$, $D = 2$ construction of Proposition 7.6, $D\lambda \approx 0.060918$. The compensator closes at $Y/F \approx 1.405161$ in the high-accuracy reconstruction, yielding $\mathcal{C}_{0.9} \approx 0.020811$. The outward-rounded certificate gives $1.4047874 < Y/F < 1.4055543$, which rigorously implies $\mathcal{C}_{0.9} > 0.020$ for the specified trajectory. The threshold $\delta = 0.020$ requires $Y/F \le \exp(9D\lambda - 10(0.020)) \approx 1.416609$, which the certified path satisfies. Thus the receiving asset's terminal price determines whether the same broad lower-limit streak qualifies as a systemic loss event.

7.3 A minimum speed for a consecutive closing streak

The second gate is *single-asset execution capacity.* A closing price at the lower limit must be reached in finite time despite recovery toward F. For $g(S) \equiv \eta$, start a session at price s and suppose its length is τ and its lower floor is $(1-\beta)s$. Even before asking who purchases the matched dollars, the gross share-rate cap bounds the fastest decline through $\dot{z} = -\rho(z - F) - \eta\bar{v}$, $z(0) = s$. This yields a necessary threshold but leaves the market-wide matching problem unresolved.

Proposition 7.2 (Rolling-limit shock-rate certificates)**.** *In the constant per-share impact model, a lower-limit close in a session starting at $s > 0$ requires*

$$\boxed{\bar{v} \ge \left[\frac{\rho}{\eta}\left(F - s + \frac{\beta s}{1 - e^{-\rho\tau}}\right)\right]_+.} \tag{7.5}$$

If the same asset closes at the same proportional lower limit on each of D sessions of common duration τ, with $s_d = Fq^{d-1}$, then

$$\boxed{\bar{v} \ge \frac{\rho F}{\eta} \max_{1 \le d \le D}\left[1 + q^{d-1}\left(\frac{1-q}{1 - e^{-\rho\tau}} - 1\right)\right].} \tag{7.6}$$

Moreover, when $\eta\bar{v} < \rho F$, all prices starting at F obey the perpetual rate-only *floor $S_i(t) \geq F - \eta\bar{v}/\rho$, implying the necessary integer-streak bound*

$$\boxed{D \leq \frac{\log(1 - \eta\bar{v}/(\rho F))}{\log q}.} \tag{7.7}$$

The bound is understood only for $\eta\bar{v} < \rho F$; with a larger rate cap it supplies no finite-day obstruction. None of these scalar inequalities is sufficient for multi-asset price-band feasibility or terminal dollar closure.

Proof. Variation of constants and $\dot{\theta}_i \geq -\bar{v}$ imply $S_i(\tau) \geq F + (s - F)e^{-\rho\tau} - (\eta\bar{v}/\rho)(1 - e^{-\rho\tau})$. Requiring $S_i(\tau) \leq (1 - \beta)s$ yields (7.5); substitute $s = Fq^{d-1}$ for each day to obtain (7.6). The same scalar comparison from initial F gives $S_i(t) \geq F - (\eta\bar{v}/\rho)(1 - e^{-\rho t})$, hence the uniform floor; combine it with Fq^D and note $\log q < 0$ to deduce (7.7). □

7.4 Why universal consecutive limit-down can be impossible

The third gate is *market-wide conservation*. The all-asset case is especially restrictive: there is no complement capable of absorbing a net dollar flow. The same potential that gave the original stability barriers now rules out a second simultaneous decline under a simple below-benchmark band condition. This is not a separate economic mechanism; it is the global end of the breadth distinction already exposed by the A/B/C regimes.

Theorem 7.3 (An obstruction to a second all-asset limit-down close)**.** *Assume $g_i(S) > 0$, $\rho_i > 0$, unchanged fundamental benchmarks $F_i > 0$, and instantaneous interasset executed-dollar balance $\sum_i \mathrm{d}X_i = 0$. Suppose every asset closes the first session at $S_i(T_1) = q_i F_i < F_i$, and the* entire second-session admissible band *lies at or below its own benchmark,*

$$(1 + \beta_{i,2,+})q_i \leq 1 \quad \textit{for every } i. \tag{7.8}$$

Then it is impossible for every *asset to close the second session strictly below its first close. In particular, simultaneous second-day lower-limit closes in all N assets are impossible. This holds regardless of the transaction-rate cap, terminal signed-dollar balances, or the functional form of the positive g_i.*

Proof. Use $\Phi_i'(S) = S/g_i(S) > 0$ and $H = \sum_i \Phi_i(S_i)$. Trade settlement preserves H, and recovery contributes $\dot{H} = -\sum_i \rho_i S_i(S_i - F_i)/g_i(S_i) \geq 0$ whenever every $S_i \leq F_i$. Condition (7.8) keeps the full second-session price path in that region. Hence $H(T_2) \geq H(T_1)$. If every second close were below the preceding close, strict monotonicity of all Φ_i would instead give $H(T_2) < H(T_1)$, a contradiction. □

For the stylized symmetric proportional band, $q_i = 1 - \beta_i$ and $\beta_{i,2,+} = \beta_i > 0$, so $(1+\beta_i)(1-\beta_i) = 1 - \beta_i^2 < 1$ and the theorem applies. The conclusion is a property of our *fully executed, cross-asset dollar-balanced, fixed-fundamental* market, not a prediction that all stocks cannot fall together again on a real exchange: unmatched orders, cash as an external trading leg, and fundamentals are absent. It also does not prohibit $k < N$ assets from having a consecutive streak while other assets supply the offsetting dollar flows.

The same model cannot represent a positive-duration *all-asset locked limit-down state* with every price fixed strictly below F_i: maintaining $\dot{S}_i = 0$ would require $\dot{X}_i = \rho_i S_i(S_i - F_i)/g_i(S_i) < 0$ for

every asset, contradicting $\sum_i \dot{X}_i = 0$. A limit-hit closing print is instead permitted as an endpoint condition. Persistent unmatched sell queues require a new order-book state, not simply an extended price trajectory.

7.5 A positive multi-day resource and accounting constraint

Even when not all assets fall, each target that is below F faces upward recovery. Terminal dollar closure therefore requires a path-history resource, not just a final selling flow: earlier above-benchmark excursions must provide the positive recovery contribution that compensates later losses. This is another expression of the same potential introduced in Section 3, now integrated over successive trading sessions.

Proposition 7.4 (Preconditioning needed for a dollar-closed limit-down streak)**.** *Let an asset begin at F and finish D sessions at $s_D < F$, with $X_i(T_D) = 0$. Then its trajectory necessarily visits $S_i > F$ on a set of positive trading-time measure and satisfies*

$$\rho \int_0^{T_D} \frac{S_i(t)(S_i(t) - F)}{g(S_i(t))} \, \mathrm{d}t = \int_{s_D}^{F} \frac{u}{g(u)} \, \mathrm{d}u > 0. \tag{7.9}$$

If every session after the first keeps the asset at or below F, and the first-session upper band is $U_1 > F$, a necessary condition is

$$\boxed{\rho\tau_1 \max_{F \le u \le U_1} \frac{u(u - F)}{g(u)} \ge \int_{s_D}^{F} \frac{u}{g(u)} \, \mathrm{d}u.} \tag{7.10}$$

For constant impact $g = \eta$, and $s_D = Fq^D$, $U_1 = F(1 + \beta_+)$, this simplifies to $\rho\tau_1 \ge (1 - q^{2D})/[2\beta_+(1 + \beta_+)]$. Writing $H_+ = \rho\tau_1\beta_+(1 + \beta_+)$, whenever $2H_+ < 1$ this imposes the additional necessary streak-length bound

$$\boxed{D \le \frac{\log(1 - 2H_+)}{2 \log q}.} \tag{7.11}$$

This bound is specific to terminal dollar closure and the absence of above-fundamental excursions in later sessions. It is not sufficient to produce a consecutive limit-down path.

Proof. Integrate the identity $\mathrm{d}\Phi(S_i) = \mathrm{d}X_i - \rho S_i(S_i - F)\,\mathrm{d}t/g(S_i)$ over the full horizon. Since $X_i(T_D) = 0$ and $\Phi(F) = 0$, the integral must equal $-\Phi(s_D) > 0$. Any portion of the trajectory at $S_i \le F$ contributes a nonpositive integrand, so the positive part must arise above F. If later sessions cannot rise above F, all positive contribution must be earned in session 1, whose supremum is bounded by the left side of (7.10). □

7.6 A necessary receiving-capacity bound for multiasset execution

The preceding tests constrain the falling stocks individually or the entire market collectively. Neither captures an intermediate case central to systemic crises: *most* assets fall, while a small complement remains to settle the opposite dollars. This is precisely the $k < N$ case behind the A/B/C construction. Structural reachability says a compensator can exist; execution feasibility asks *how much its allowed gross share rate and price band can absorb.* We partition the assets into a falling subset J and its complement J^c. Conservation and per-asset gross rates directly bound how many executed dollars can cross this partition. This is a *necessary subset-capacity inequality*

for a complete trading graph with node-level gross-rate caps and no separate edge capacities. The inequality connects a structurally permitted partial-coverage loss to the ability of the non-falling assets to receive the required executed dollars.

Proposition 7.5 (A market-cut capacity and consecutive-limit lower bound)**.** *Fix a subset J of assets and write J^c for its complement. Under paired-dollar execution and the per-asset gross share-rate cap $\bar{v}$, every admissible absolutely continuous path satisfies, for almost every t,*

$$\boxed{\left|\sum_{i\in J}\dot{X}_i(t)\right| \le \bar{v}\min\left\{\sum_{i\in J}S_i(t),\ \sum_{j\in J^c}S_j(t)\right\}.} \tag{7.12}$$

Suppose that throughout a session $[T_{d-1},T_d]$ of length τ_d, each $i\in J$ stays at or below its unchanged benchmark F_i and ends at $s_{i,d}<s_{i,d-1}$. Let $\Phi_i'(S)=S/g_i(S)>0$ and put

$$K_{J,d}:=\sum_{i\in J}\left[\Phi_i(s_{i,d-1})-\Phi_i(s_{i,d})\right]>0. \tag{7.13}$$

If all compensating assets satisfy $S_j(t)\le U_{j,d}$ during that session, then the session's paired executed-dollar turnover M_d and the required gross-share-rate cap satisfy

$$\boxed{M_d\ge K_{J,d},\qquad \bar{v}\ge\frac{K_{J,d}}{\tau_d\sum_{j\in J^c}U_{j,d}}.} \tag{7.14}$$

When $J^c=\varnothing$, the hypothesized simultaneous strict decrease of all assets cannot occur on a session lying entirely below benchmarks.

Proof. Dollar pairing gives $\sum_{i\in J}\dot{X}_i=-\sum_{j\in J^c}\dot{X}_j$. In either set, $|\dot{X}_i|\le S_i r_i^{\mathrm{gross}}\le S_i\bar{v}$, proving (7.12). The price-response coordinate obeys $\Delta X_i=\Phi_i(s_{i,d})-\Phi_i(s_{i,d-1})+\rho_i\int_{T_{d-1}}^{T_d}S_i(S_i-F_i)/g_i(S_i)\,\mathrm{d}t$. The integral is nonpositive for $i\in J$, hence $\sum_{i\in J}\Delta X_i\le -K_{J,d}$. The net executed dollars entering J^c are therefore at least $K_{J,d}$. The total matched-dollar turnover must cover this flow, and the complement's gross-share capacity is no more than $\bar{v}\tau_d\sum_{j\in J^c}U_{j,d}$. These yield (7.14). If J^c is empty, there is no opposing flow. □

Notice how the structural and execution arguments now meet. If J^c is empty, the cut has no receiving side, echoing the all-asset obstruction. If J^c is nonempty, a path may be structurally permitted but still require a rate above the complement's available capacity. The same remaining assets have a second, opposing effect on the crisis functional: their price appreciation can offset the falling majority's average log loss, as measured by Proposition 7.1. Feasibility therefore requires *both* sufficient funding absorption and sufficient uncompensated average loss.

For constant per-share impact $g_i\equiv\eta$, $k=|J|$ identical lower-limit assets, initial benchmark F, common lower-limit ratio $q\in(0,1)$, and $m=N-k$ compensators with common upper-band fraction $\beta_{c,+}$, the dth-day specialization is

$$\boxed{\bar{v}\ge\frac{k}{m}\,\frac{F\left(q^{2(d-1)}-q^{2d}\right)}{2\eta\tau_d(1+\beta_{c,+})^d}\quad(d\ge 2),} \tag{7.15}$$

provided all k target prices remain below F and each compensator's preceding reference price has evolved from F under the specified upper bands. The k/m multiplier makes the funding bottleneck explicit. This is a *necessary* condition, not an exact optimum. For the ten-asset, two-day example ($k=9,m=1,q=0.97,\beta_{c,+}=0.25,\tau_d=20$), it gives only $\bar{v}\gtrsim 0.0080$; the stronger systemic quadratic bound below is more informative for that particular configuration.

7.7 Validated lower-rate construction and a two-sided minimum-rate bracket

Our necessary constraints do not establish simultaneous feasibility. We now certify a *single path* obeying all of them. Fix $N = 10$, with nine identical target assets and one compensator, $F = \eta = 1$, $\rho = 0.4$, two sessions of length 20, target rolling bands $\pm 3\%$, compensator bands $\pm 25\%$, and terminal closure $X_i(40) = 0$ for every asset. The targets share one prescribed continuous price path $x(t)$; the compensator follows the scalar nonlinear ODE determined by paired dollars. Hence the construction reduces to two price states, but certifies a nonempty admissible set in the full ten-asset model, not generic robustness to heterogeneous parameters. The following piecewise-linear target path admits an independently rerunnable fixed-point interval certificate of its nonlinear compensator dynamics and execution constraints.

Proposition 7.6 (A computer-assisted two-day rate certificate)**.** *In the ten-asset specification just stated in Section 7.7 ($k = 9$, $F = \eta = 1$, $\rho = 0.4$, $\tau_1 = \tau_2 = 20$, target bands $\pm 3\%$, compensator bands $\pm 25\%$), there exists a continuous, absolutely continuous, fully paired-dollar path with all per-asset terminal signed-dollar ledgers exactly zero such that the same nine assets close at* 0.97 *and* 0.9409*, respectively, while*

$$\boxed{\mathcal{C}_{0.9}(40) > 0.020, \quad \bar{v}_{\text{used}} < 0.285, \quad M < 3, \quad H_i^- < 2 \quad (1 \le i \le 10).} \tag{7.16}$$

Every executed price remains inside its rolling session band. Here M counts paired executed dollars once and H_i^- denotes gross shares sold, so initial holdings of two shares of every asset are sufficient without shorting.

Proof outline; reproducible validated enclosure in Appendix G. Let $r = 23/10$. Prescribe the common target price $x(t)$ as the continuous piecewise-linear interpolation of

$$(t, x) = (0, 1), (r, h), (20 - r, h), (20, 0.97), (20 + r, 0.995), (40 - r, 0.995), (40, 0.9409). \tag{7.17}$$

The first plateau h is the unique root near 1.02 of the exact net-dollar-closure equation

$$1016000000000h^2 - 97138000000h - 6646157699 = 0, \quad 1.02020225507905 < h < 1.02020225507908. \tag{7.18}$$

The compensator $y(t)$ is the unique positive solution of $\dot{y} = -0.4(y - 1) - 9u(t)/y$, $y(0) = 1$, where $u(t) = x(t)[\dot{x}(t) + 0.4(x(t) - 1)]$. All ten executed-dollar flows sum to zero at each time, and the exact root equation gives $X_i(40) = 0$ for every target, hence also for the compensator. A validated interval integration gives

$$\begin{gathered} 1.2219388 < y(20) < 1.2225356, \quad 1.4047874 < y(40) < 1.4055543, \\ \sup_t 9|u(t)|/y(t) < 0.284481, \\ \int_0^{40} 9[u(t)]_+/y(t)\,\mathrm{d}t < 1.880, \qquad 9\int_0^{40} |u(t)|\,\mathrm{d}t < 2.984. \end{gathered} \tag{7.19}$$

The targets have even smaller gross rates and gross selling shares; sessionwise interval enclosures verify both rolling bands. Exact rational enclosures for $e^{0.2}$ verify $\overline{L}(40) > 0.020$, while the nine target losses are $-\log(0.9409) > 0.020$, proving (7.16). □

For this *fixed* ten-asset, two-day, nine-streaking-assets instance, let $\bar{v}_*$ be the infimum gross rate cap among *all* admissible paths also satisfying $\mathcal{C}_{0.9}(40) \ge 0.020$. The quadratic system-wide

bound (6.11) implies for every $\bar{v} \le 0.044$ (with $F = \eta = 1$, $\rho = 0.4$, $T = 40$):

$$\mathcal{C}_{0.9}(40) \le (\bar{v}/0.4)^2 \left[1 + \frac{1}{2(1 - \bar{v}/0.4)^2}\right] \le \frac{1563441}{79210000} < 0.020.$$

Since the right-hand side of (6.11) is continuous in $\bar{v}$ near 0.044 and is strictly below 0.020 at that value, the exclusion extends to some rates above 0.044. Together with the feasible construction in Proposition 7.6, this yields the rigorous, though nonoptimal, bracket for the minimum gross-rate cap:

$$\boxed{0.044 < \bar{v}_* < 0.285.} \tag{7.20}$$

This bracket completes the paper's chain of questions for one fixed market: the response law permits partial-coverage crashes; the crisis definition requires sufficient average loss; daily price limits and gross rates restrict the admissible paths; paired-dollar settlement constrains the receiving side; and a certified path actually reaches the target. The lower bound and upper witness serve different logical roles. The certified control supplies a feasible upper rate, while the quadratic barrier supplies a strict lower bound. Determining the exact global minimum remains a constrained multiday optimal-control problem.

Sensitivity and the role of symmetry. The ten-asset certificate deliberately reduces nine identical target states to one trajectory. Holding that *same* $x(t)$ and exact target-dollar closure fixed, a numerical sensitivity check repartitions ten assets into k identical targets and $10 - k$ identical compensators. Their common compensator price obeys $\dot{y} = -0.4(y - 1) - [k/(10 - k)]u(t)/y$. Table 2 shows why the number of receiving assets matters: the specified control achieves the fixed 0.020 mean-loss target only for $k = 9$ among the cases tested. Changing k also changes the coverage fraction $k/10$; it is not a test at fixed $\alpha = 0.9$. This is *not* a proof that no redesigned path works for $k < 9$; those rows are high-accuracy ODE diagnostics, unlike the $k = 9$ interval certificate. Similarly, for the fixed $k = 9$ path, the certificate proves compliance at a symmetric compensator band of $\pm 24.7\%$ and proves violation at $\pm 24\%$ (the price interval at $t = 17.7$ lies strictly below 0.76). Changing the band does not rerun or optimize the control.

Table 2: Sensitivity of the *fixed certified target-control shape* to cross-sectional composition. Only the $k = 9$ row is interval-certified; other rows use independently integrated, high-accuracy scalar ODEs. A common compensator band is the smallest symmetric daily fraction compatible with the *computed* path; the values are diagnostics, not optimized feasibility thresholds.

Targets k	Compensators	$y(40)$	Mean log loss	Computed band needed
6	4	1.08112	0.00535	4.28%
7	3	1.12302	0.00784	6.55%
8	2	1.20141	0.01204	10.88%
9	1	1.40516	0.02081	24.55%

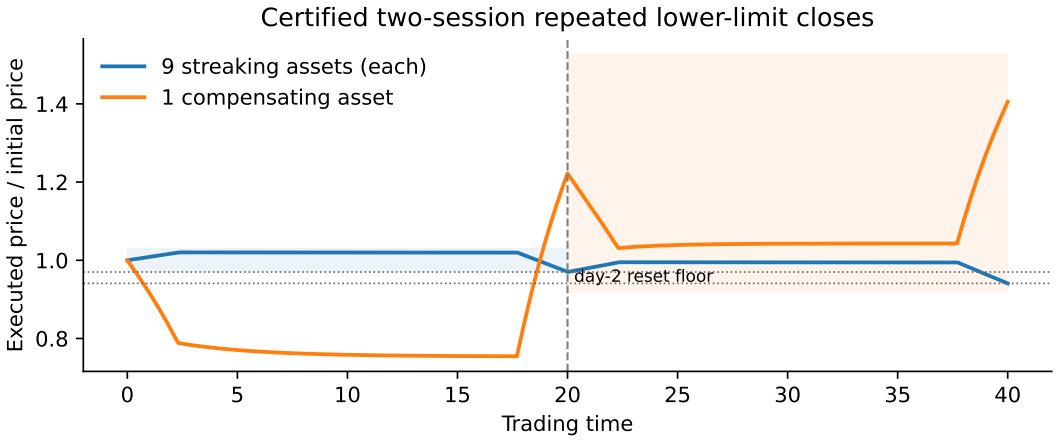


Figure 5: Interval-certified lower-rate path. Continuous target and compensator prices for the improved two-day construction. The nine targets close at the reset lower limits; the compensator stays inside its 25% daily limits. Curves are illustrative numerical reconstructions; admissibility is independently certified by the outward-rounded arithmetic described in Appendix G.

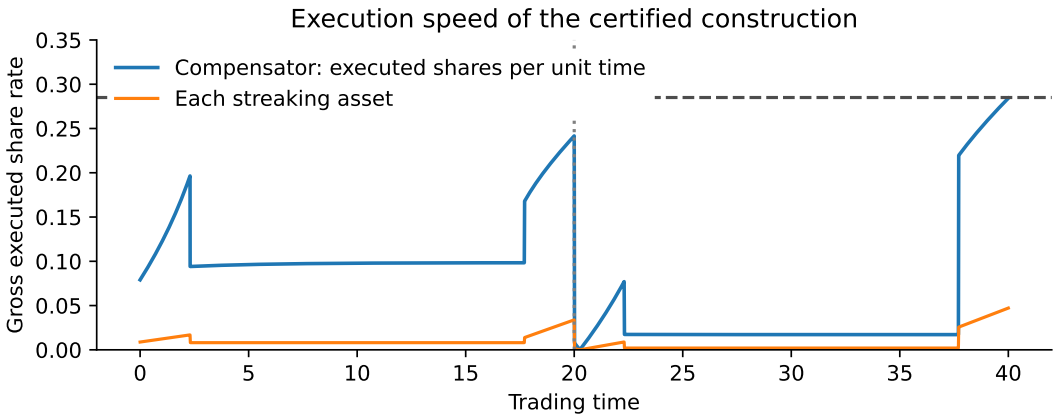


Figure 6: Gross execution rates in the certified path. The largest computed gross share-rate is approximately 0.28422; rigorous interval rounding certifies a bound strictly below 0.285, without instantaneous pulses. The cap is per asset, before netting buys and sells.

8 Interpretation, limitations, and testable implications

The generalized-coordinate normal form (2.6) organizes the results into one sequence. Executed dollars enter additively and conserve the sum of response-adapted coordinates; recovery drives that sum between trades. Curvature yields global stability barriers, the A/B/C family shows how the barriers change with response geometry, and coordinate distances bound the execution resources required to cross into a systemic loss set. Finite rates and resetting daily bands further restrict admissible paths. The certified two-day construction demonstrates that these restrictions can be met simultaneously, while the subset-capacity inequality links the falling majority to the buying capacity of the remaining assets.

The distinction from equilibrium-driven mechanisms is precise. Huang and Wang [10] study nonsynchronous demand, Cespa and Foucault [4] cross-asset learning, Cont and Wagalath [6] distressed-sale induced correlations, and Cordoni and Lillo [7] strategic instability in impact games. We abstract from those selection mechanisms and instead identify conditions on a specified *controlled market* that exclude a crisis state or permit a constrained path to it. Kyle and Obizhaeva [15] quantify the price implications of large trades and execution speed using invariance and observed crashes; our capacity and lower-rate bounds make no calibrated quantitative prediction for observed crashes. The contribution is the joint structural–resource–execution reachability question under common settlement and recovery rules, rather than a substitute for the existing mechanisms.

The ten-asset witness uses a symmetric nine-target trajectory and one compensator, reducing computation to two price states. The sensitivity checks in Table 2 quantify how the compensator burden changes under a fixed target control. The certified event is a *systemic drawdown*: nine stocks each lose 5.91% across two sessions while the tenth appreciates, and the equal-weight geometric mean falls by roughly 2.06%.

Empirical application would require identifying $g(S)$, recovery ρ, executable cross-asset order flow, inventories and cash settlement from high-frequency data. A static order-book snapshot or daily OHLCV series is insufficient to identify these inputs; the effective response g is not automatically a Kyle lambda. Exchange bands in the construction are hypothetical percentage restrictions on *executed prices*, whereas actual Tokyo Stock Exchange limits usually use price-tiered absolute yen ranges and specified enlargement rules [12]. A modeled lower-limit close does not establish an unfilled sell queue or a genuine locked-limit market.

Dollar-flat reallocation can change final shareholdings and dissipate trading wealth even though isolated share-flat round trips have nonnegative executed cost. The results characterize admissible price states of this controlled market; behavioral selection and empirical event frequencies require further models.

Diagonal impact and potential-based extensions. Our cross-asset dependence is generated by executed-dollar pairing, while the marginal impact matrix is diagonal. In models with direct cross-impact $\mathrm{d}\boldsymbol{S} = \mathsf{G}(\boldsymbol{S})\,\mathrm{d}\boldsymbol{\theta}$ and invertible G, the transaction accounting suggests the one-form $\omega = \mathbf{1}^{\top}\operatorname{diag}(\boldsymbol{S})\mathsf{G}(\boldsymbol{S})^{-1}\,\mathrm{d}\boldsymbol{S}$. A scalar transaction potential would require appropriate exactness conditions on this one-form and a compatible sign condition for the recovery drift. Symmetry or positive definiteness of a matrix impact coefficient alone does not establish these properties. This distinguishes the present separable-potential results from matrix-kernel conditions studied in Alfonsi, Klöck, and Schied [1] and the cross-impact Nash games of Cordoni and Lillo [7].

9 Conclusion

The response-adapted generalized price coordinate reduces nonlinear trading impact to additive, dollar-balanced transfers and a restoring drift. Its curvature identifies global exclusion barriers; its coordinate distances yield necessary resource bounds. A common response family exhibits three reachability regimes, and a bounded-rate, multiday interval-certified construction shows that a systemic drawdown can survive the tighter execution restrictions. Structural feasibility, available resources, and multiday execution capacity thus form successive, mathematically distinct tests of market stability.

Code and numerical data availability

The accompanying reproducibility package includes `verification/certify_streak_interval.py` (exact rational root isolation and outward-rounded fixed-point enclosures), `verification/verify_consistency.py`, numerical sensitivity scripts, figure-generating code, machine-readable outputs, the LaTeX source, and a file-hash manifest. Reproduction commands and software requirements are in `REPRODUCE.md`. The certificate was rerun from source, with independent floating-point ODE and exact-rational consistency checks. A versioned, checksum-indexed reproducibility archive accompanies this manuscript; the computations are tied to the supplied source and certificate versions.

Declaration on AI-assisted tools

AI-assisted tools were used in drafting and organizing parts of the manuscript, developing and checking mathematical calculations, writing computational scripts, and preparing figures. The author assumes responsibility for all content, including mathematical claims, source attribution, and the reproducibility of the numerical certificate.

A Settlement, coordinate changes, and variation bounds

The continuous response rule is interpreted along finite-variation execution paths. During an idealized pulse, physical time is frozen ($\mathrm{d}t = 0$) and the states move along the response ODE as the executing share size varies. The pulse's executed dollar transfer is the *integral* $\int S_i \,\mathrm{d}\theta_i$, not the beginning-of-pulse price times the final share increment. By continuity of $g > 0$, its endpoints can be joined by an admissible response path whenever the intervening price stays positive.

For one asset the differential rule reads

$$\mathrm{d}\theta = \frac{\mathrm{d}S + \rho(S - F)\,\mathrm{d}t}{g(S)}. \tag{A.1}$$

Decompose $S = F + (S - F)$ and integrate:

$$\int_0^T S \,\mathrm{d}\theta = F[\theta(T) - \theta(0)] + \int_F^{S(T)} \frac{u - F}{g(u)}\,\mathrm{d}u + \rho \int_0^T \frac{(S - F)^2}{g(S)}\,\mathrm{d}t,$$

where the initial price is F. This proves (2.11). For $S < F$ the primitive $E(S)$ in (2.10) is also nonnegative because the integration orientation reverses the sign of the negative integrand. Hence if $\theta(T) = \theta(0)$, the total executed cost cannot be negative.

Let $y = \Phi(S)$ and $z = \Psi(S)$, with

$$\Phi'(S) = \frac{S}{g(S)}, \qquad \Psi'(S) = \frac{1}{g(S)}.$$

Then

$$\mathrm{d}y = \mathrm{d}X - \rho\frac{S(S-F)}{g(S)}\,\mathrm{d}t, \qquad \mathrm{d}z = \mathrm{d}\theta - \rho\frac{S-F}{g(S)}\,\mathrm{d}t. \tag{A.2}$$

The recovery terms have signs opposite to $S - F$. Thus whenever $y < 0$ (equivalently $S < F$), the no-trade drift in y is positive; likewise for z. By separating the positive and negative variations in (A.2), a trajectory ending at $y(T) \le \Phi(s) < 0$ must accumulate at least $-\Phi(s)$ in negative executed-dollar variation. One way to see this is to start at the *last* time when y equals zero (or $t = 0$ if no later zero exists). On the remaining interval $y \le 0$ and the recovery drift is nonnegative. The negative transaction variation must therefore account for the entire fall to $y(T)$. The same argument in z gives the share-sale bound. Summing sales across qualifying assets counts each paired-dollar transfer exactly once. If a terminally qualifying asset has inventory cap H_i and $X_i(T) = 0$, equation (2.11) implies

$$E(S_i(T)) \le F\,[\theta_i(0) - \theta_i(T)] \le FH_i.$$

Since $E(S)$ increases as S falls below F, $S_i(T) \le s$ implies $E(s) \le FH_i$. This proves Theorem 5.1.

When $\rho = 0$, $\mathrm{d}\Phi(S_i) = \mathrm{d}X_i$. Terminal dollar closure therefore implies $\Phi(S_i(T)) = \Phi(F) = 0$ and, by strict monotonicity, $S_i(T) = F$. Within the present separable response class, recovery makes the nontrivial dollar-closed endpoint effect possible.

B Local two-pulse and three-pulse expansions

We document the expansion used in Propositions 4.3 and 4.2. Normalize local variations by $x = S/F - 1$ and $c = g(F)/F^2$. At $S = F$, the executed-dollar coordinate $y = \Phi(S)$ satisfies $\Phi(F) = 0$, $\Phi'(F) = F/g(F)$, and

$$\frac{\mathrm{d}S}{\mathrm{d}y} = \frac{g(S)}{S}, \qquad \frac{\mathrm{d}^2S}{\mathrm{d}y^2} = \frac{g(S)[Sg'(S) - g(S)]}{S^3}. \tag{B.1}$$

Hence, writing $e_0 = e(F)$, the price ratio after a small dollar-coordinate displacement y is

$$\frac{S}{F} = 1 + cy + \frac{1}{2}(e_0 - 1)c^2y^2 + O(y^3), \tag{B.2}$$

with constants uniformly bounded for paths in a compact neighborhood of F.

For the two-pulse experiment, the first impulse displaces y by hu_i at time zero. It produces the initial price response according to (B.2). Recovery over T multiplies the *level* displacement $S - F$ by $r = e^{-\rho T}$; the terminal pulse subtracts the same dollar amount hu_i from the executed-dollar coordinate. Expanding Φ and Φ^{-1} around F gives

$$L_i(T) = c(1-r)hu_i + \frac{c^2h^2}{2}\,\lambda(r, e_0)u_i^2 + O(h^3), \tag{B.3}$$

$$\frac{1}{N}\sum_i L_i(T) = \frac{c^2h^2}{2N}(1-r)(2-e_0-r)\sum_i u_i^2 + O(h^3),$$

where the individual second-order coefficient λ is not needed separately. The linear term vanishes in the mean because $\sum_i u_i = 0$. If at least k components are positive, those assets have positive first-order losses and the kth largest loss is positive for small h; the mean is positive under the stated time-curvature condition.

For the three-pulse experiment, the dollar-coordinate impulses at $0, T/2, T$ have amounts hu_i, $-(1+r)hu_i$, and rhu_i with $r = e^{-\rho T/2}$. Direct substitution and second-order expansion yield

$$L_i(T) = \frac{1-e_0}{2} c^2 h^2 u_i^2 r(1-r^2) + O(h^3), \tag{B.4}$$

with first-order terms canceling exactly. If $e_0 < 1$ and every $u_i \neq 0$, all terminal losses are positive for sufficiently small h. Since every pulse amount vector is zero-sum and the amounts cancel componentwise over the full horizon, the construction also satisfies the paired-dollar and terminal dollar-closure rules. The initial inventory requirement and turnover are finite for fixed h.

C Uniform-resource severe drawdowns for partial coverage

We prove Theorem 4.1(C) constructively. Fix $N \geq 2$, set $m = N-1$, and choose the constant-response model $g(S) = \eta > 0$. We may assume the target crisis threshold δ is large, since a path exceeding a larger threshold automatically exceeds every smaller one. It is enough to produce m terminal prices $\varepsilon \downarrow 0$ with a final "compensator" asset, because every $k < N$ satisfies $k \leq m$.

C.1 A radial representation of paired-dollar trading

For $g \equiv \eta$,

$$\mathrm{d}(S_i^2) = -2\rho S_i(S_i - F)\,\mathrm{d}t + 2\eta\,\mathrm{d}X_i, \qquad \mathbf{1}^\top \mathrm{d}\boldsymbol{X} = 0. \tag{C.1}$$

Let $R = (\sum_i S_i^2)^{1/2}$. A paired-dollar trade keeps R fixed; a positive-price change at fixed R is implementable as a finite-variation path in the positive orthant of that sphere. In particular, interpolate the squared prices $Y_i = S_i^2$ linearly between two positive vectors of the same sum: every intermediate $Y_i > 0$ and $\sum_i \mathrm{d}Y_i = 0$, so $\mathrm{d}X_i = \mathrm{d}Y_i/(2\eta)$ is a legal paired transfer. The corresponding share increment is $\mathrm{d}\theta_i = \mathrm{d}S_i/\eta$.

Select a small constant $c > 0$ and a parameter $b \in (0, \beta)$, where one may take $\beta = 1/5$, so that

$$a = \sqrt{1 - b^2 - (m-1)c^2} > b > c > 0, \quad H = a + (m-1)c + b, \quad bH < 1. \tag{C.2}$$

Such choices exist for sufficiently small c. During stage $j \in \{1, \dots, m\}$, fix the spherical direction

$$\boldsymbol{u}_j = (c, \dots, \underbrace{a}_{j}, \dots, c, b), \qquad \|\boldsymbol{u}_j\|_2 = 1. \tag{C.3}$$

Set $\boldsymbol{S}(t) = R(t)\boldsymbol{u}_j$ throughout the stage. To make it solve the price equations, use

$$R(t) = F\left[H + (\sqrt{N} - H)e^{-\rho t}\right], \qquad \dot{\theta}_i = \frac{\rho F}{\eta}\left(u_{j,i}H - 1\right). \tag{C.4}$$

Indeed, differentiating $S_i = u_{j,i}R(t)$ in the price equation gives (C.4) and

$$\dot{X}_i = \frac{\rho F R(t)}{\eta}\, u_{j,i}(u_{j,i}H - 1), \quad \sum_i \dot{X}_i = \frac{\rho F R(t)}{\eta}\left(H\sum_i u_{j,i}^2 - \sum_i u_{j,i}\right) = 0, \tag{C.5}$$

because $\sum_i u_{j,i}^2 = 1$ and $\sum_i u_{j,i} = H$.

C.2 Closing the compensator's dollar flow

At $t = 0$ begin with $F\mathbf{1}$ and apply a positive-price pulse at fixed radius $F\sqrt{N}$ to the first direction $\boldsymbol{u}_1$. Switch between stages $\boldsymbol{u}_j$ using similar fixed-radius pulses. At time T use a final fixed-radius pulse to reach

$$S_1(T) = \cdots = S_m(T) = \varepsilon, \quad S_N(T) = \sqrt{R(T)^2 - m\varepsilon^2}. \tag{C.6}$$

The terminal price vector is positive for all sufficiently small ε.

For constant g, integrating (C.1) gives the exact executed-dollar ledger

$$X_i(T) = \frac{S_i(T)^2 - F^2}{2\eta} + \frac{\rho}{\eta}\int_0^T S_i(t)[S_i(t) - F]\,\mathrm{d}t. \tag{C.7}$$

For the compensator, its stage weight is always b. Define

$$\mathcal{F}_\varepsilon(b) = \frac{R_b(T)^2 - m\varepsilon^2 - F^2}{2\eta} + \frac{\rho}{\eta}\int_0^T bR_b(t)[bR_b(t) - F]\,\mathrm{d}t, \tag{C.8}$$

where R_b depends continuously on b through (C.4). For sufficiently small c and ε, $\mathcal{F}_\varepsilon(0) > 0$ because $R_0(T) > F$. At $b = \beta$, (C.2) gives $bH < 1$, so the time integral is negative at a rate proportional to T, while the terminal term stays bounded. Thus for some fixed, sufficiently long T one has $\mathcal{F}_\varepsilon(\beta) < 0$, uniformly for all small ε. By continuity there is $b(\varepsilon) \in (0, \beta)$ for which $X_N(T) = 0$ exactly.

C.3 Closing every target asset's dollar flow

It remains to make the other m terminal dollar balances equal. Whenever an asset is in the high state a rather than the low state c, the difference in its recovery-integrand is

$$W(t) = (a - c)R(t)\big[(a + c)R(t) - F\big]. \tag{C.9}$$

This is strictly positive: $(a + c)H = 1 + b(a - b) + mac + bc > 1$, and $R(t) \geq FH$. Consequently $W(t) > 0$ for $t \in [0, T]$. Choose a partition $0 = t_0 < \cdots < t_m = T$ such that

$$\int_{t_{j-1}}^{t_j} W(t)\,\mathrm{d}t = \frac{1}{m}\int_0^T W(t)\,\mathrm{d}t, \qquad j = 1, \ldots, m. \tag{C.10}$$

The continuous, strictly increasing cumulative integral makes this partition possible. Let target asset j occupy the high state in stage j and the low state otherwise. All target assets have the same initial price F, terminal price ε, and recovery integral in (C.7), hence identical terminal balances $X_i(T)$. Every pulse and every stage respects $\sum_i \mathrm{d}X_i = 0$; since $X_N(T) = 0$, the common target balance must also vanish. Therefore

$$X_i(T) = 0 \qquad \text{for every } i = 1, \ldots, N. \tag{C.11}$$

C.4 Uniform finite turnover and inventory

The path remains in the positive sphere of radius at most $F\sqrt{N}$, so $0 < S_i \leq F\sqrt{N}$ throughout. There are $m + 1 = N$ fixed-radius pulses, and each pulse can be implemented by monotone componentwise interpolation in squared-price coordinates. Consequently the cumulative paired-dollar pulse turnover has an explicit, though loose, upper bound depending only on N, F, η. For

example, a safe bound is $N^3F^2/(4\eta)$. During recovery-controlled stages, $\sum_i |\dot{X}_i|/2 \le \rho N F^2/\eta$, giving

$$M(\gamma) \le M_0 := \frac{N^3F^2}{4\eta} + \frac{\rho N F^2 T}{\eta}. \tag{C.12}$$

Likewise, share sales over N pulses plus the continuous stages are bounded by

$$H_i^- \le H_0 := \frac{NF\sqrt{N} + \rho FT(\sqrt{N}+1)}{\eta}, \tag{C.13}$$

independently of ε. Starting with at least H_0 shares of each asset rules out short sales. The compensator's price stays bounded and positive; its log loss remains finite. The other m losses are $\log(F/\varepsilon) \to \infty$. Hence their kth order statistic and the average loss both diverge when $k \le m$. Thus $\mathcal{C}_\alpha(T) \to \infty$ with fixed finite M_0, H_0, proving Theorem 4.1(C). □

D Construction proving the sharp full-coverage bound

The necessary radial inequality was established in the proof of Theorem 4.4. Here is the sufficiency construction. Fix $0 < b < 1/\sqrt{N}$, and choose

$$a = \sqrt{1-(N-1)b^2} > b, \quad H = a + (N-1)b, \quad R_b(t) = F\left[H + (\sqrt{N} - H)e^{-\rho t}\right]. \tag{D.1}$$

Use the N positive unit-vector directions given by permutations of $(a, b, \dots, b)$. At time zero pulse from $F\mathbf{1}$ to $R_b(0)\boldsymbol{u}_1$, then follow $S_i(t) = u_{j,i}R_b(t)$ using the same paired controls as in (C.4). When changing directions, use finite fixed-radius positive squared-price pulses. At time T, make a final pulse to the equal-price point

$$S_i(T) = \frac{R_b(T)}{\sqrt{N}}, \qquad i = 1, \dots, N. \tag{D.2}$$

For the recovery integrals to be identical for every asset, define

$$W_b(t) = (a-b)R_b(t)\big[(a+b)R_b(t) - F\big]. \tag{D.3}$$

Here $(a+b)H = 1 + Nab > 1$, so $W_b(t) > 0$ for $t \in [0, T]$. Split the interval into N parts with equal integrals of W_b. In each part a different asset is the sole high-price coordinate. All N assets now have the same initial price, the same terminal price, and the same recovery integral in (C.7). Their final dollar balances are therefore identical; since all trades are paired, their sum is zero and each balance is exactly zero.

As $b \downarrow 0$, the terminal equal price in (D.2) tends to

$$\frac{F}{\sqrt{N}}\left[1 + (\sqrt{N}-1)e^{-\rho T}\right],$$

whereas $b \uparrow 1/\sqrt{N}$ yields $S_i(T) \to F$. Continuity in b implies that every strictly intermediate common terminal price is reachable by a positive-price, finite-turnover, fully dollar-closed path. The corresponding log-loss range is precisely $0 < \delta < \delta_{\text{crit}}(N, T)$, completing the sharpness proof. The endpoint δ_{crit} itself would require the limiting direction with zero components and is excluded by strict price positivity. □

E Bounded-rate control: analytic details and audit protocol

For the three-bin control (6.14), the identities $\int f = 0$ and $\int e^{-\rho(T-t)} f(t)\,\mathrm{d}t = 0$ follow from equal bin widths and the geometric ratio $r = e^{-\rho T/3}$. Write $h_0 = g(F)/F$ and $h_1 = g(F)(e_0 - 1)/F^2$. The perturbation equations are $y' = -\rho y + h_0 f$, $z' = -\rho z + h_1 f y$. Since $y(0) = y(T) = 0$ and $(e^{-\rho(T-t)})' = \rho e^{-\rho(T-t)}$, one integration by parts gives the exact identity

$$\int_0^T e^{-\rho(T-t)} f(t) y(t)\,\mathrm{d}t = \frac{\rho}{2h_0} \int_0^T e^{-\rho(T-t)} y(t)^2\,\mathrm{d}t. \tag{E.1}$$

The resulting coefficient is $\rho/2$. Substituting $h_1/h_0 = (e_0 - 1)/F$ and $L_i = -\log(S_i/F)$ proves (6.15). At $p = 1$ the quadratic coefficient vanishes; this is consistent with the all-asset barrier but does not exclude a partial-coverage loss under a different bounded signal. The independent verifier integrates the nonlinear ODE directly with the dollar-rate signal, computes gross share execution before netting, and compares the finite-amplitude loss with (6.15).

A high-precision reconstruction of the severe pulse witness confirms its terminal-dollar accounting. Bounded-rate controls constitute a more restrictive admissible class, whose local construction and barrier are established separately.

F Numerical experiments and reproducibility

This appendix documents deterministic *constructed* price paths under the exact marginal-price settlement protocol. It does not simulate order arrivals, market-maker decisions, or empirical stock data. For the severe-crash witness we use $g(S) = 1$, $F = \eta = \rho = 1$. The code selects a positive price direction and numerically solves for the compensator direction and time partitions; between switches the radius and ledger have analytic primitives. At every transaction time the paired executed-dollar sum is zero, and at maturity *each* signed-dollar ledger closes to numerical accuracy.

Table 3: Selected severe partial-coverage witnesses. The turnover shown is a feasible execution resource, not the unknown minimum. All rows use target log loss $L = 5$ and parameters normalized to unity.

N	k	T	$\mathcal{C}_\alpha(T)$	Turnover M	Max. sale shares
3	2	3	3.3168	2.6179	3.9320
12	11	4	4.5781	38.6861	8.0804
20	19	5	4.7477	104.1750	10.1906

For $N = 3$, $k = 2$, $T = 3$, the same construction yields

Target log loss L	Crash functional	Feasible turnover	Maximum share sales
3	1.9845	2.6128	3.9365
5	3.3168	2.6179	3.9320
10	6.6501	2.6179	3.9319
20	13.3168	2.6179	3.9319

The plotted experiment uses a finite, high-resolution sampling grid only to display continuous recovery segments; the underlying segment values and cash-flow integrals are evaluated using analytic expressions, and terminal closure is checked separately. For the $N = 12$, $L = 5$ witness, the maximum terminal dollar-ledger residual under double precision is less than 10^{-13} in the

distributed script. Numerical tolerance is evidence of correct accounting implementation, not a substitute for the constructive proof in Appendix C.

Generate the bounded-rate figures with `verification/finite_rate_and_bands.py`; the two-day certificate with `verification/certify_streak_interval.py`; and numerical paths with `scripts/improved_multiday_witness.py`. Figures can be rebuilt with `scripts/make_figures.py` and `scripts/make_multiday_figures.py`; the checks are `verification/verify_consistency.py` and `scripts/verify_results.py`. All panels use model-generated inputs, not market observations.

Compile the paper by running `pdflatex` twice in the repository root.

G A reproducible interval certificate for the improved two-day witness

This appendix specifies a computer-assisted proof in a form that can be independently audited. The code `verification/certify_streak_interval.py` uses Python integers, rational arithmetic, and a fixed-point interval format with unit $Q = 10^{32}$; every multiplication and division rounds outwards. The rational endpoint bracket for h in (7.18) has strictly opposite polynomial signs and the polynomial is strictly increasing on that bracket, so it contains the unique relevant exact root.

On each linear target segment, write $x(t) = a + vt$ in local segment time and $u(t) = x(t)[v + \rho(x(t) - 1)]$. The compensator equation is the scalar nonautonomous ODE $\dot{y} = f(t, y) := -\rho(y - 1) - 9u(t)/y$. Its derivatives away from the finite set of target slope-change times are

$$f_y = -\rho + \frac{9u}{y^2}, \qquad f_t = -\frac{9u'}{y}, \qquad y'' = f_t + f_y f, \qquad u' = v[v + \rho(2x - 1)].$$

The exact target plateaus and slopes yield conservative whole-segment enclosures $|u| < 0.107$ and $|u'| < 0.041$. On the bootstrap strip $0.70 \le y \le 1.60$, these imply $|f| < 1.70$ and

$$|y''| \le \frac{9|u'|}{0.70} + \left(\rho + \frac{9|u|}{0.70^2}\right)\left(0.60\rho + \frac{9|u|}{0.70}\right) < 5.$$

A tighter value of this last bound is recomputed with outward-rounded rational arithmetic on each of the six target segments. On a fixed grid of $h_0 = 1/2000$, a scalar comparison principle is combined with the one-step Taylor enclosure

$$y(t + h_0) \in y(t) + h_0 f(t, y(t)) + [-B_j h_0^2/2, B_j h_0^2/2],$$

where B_j is the certified segment remainder bound. If the beginning state belongs to $[l, u]$, comparison of scalar ODE solutions gives the next lower bound from $l + h_0 f(t, l) - B_j h_0^2/2$ and the upper bound from $u + h_0 f(t, u) + B_j h_0^2/2$, each rounded outwards. A local $|f| < 1.70$ excursion bound encloses the solution between grid points and verifies, inductively, that it stays in the bootstrap strip. The bound is reinitialized at each piecewise-linear slope discontinuity without taking an artificial step across it.

There are exactly $80,000$ certified grid steps. For the improved path the machine-checkable output gives $0.7534859 < y(t) < 1.223386$ over the first day and $1.0301664 < y(t) < 1.406405$ over the second. It also encloses $1.2219388 < y(20) < 1.2225356$ and $1.4047874 < y(40) < 1.4055543$, establishing the moving 25% second-session limits and a terminal interval containing the high-precision reference value 1.405160822. For the nine target assets, the target coordinates follow

exact line segments, with their daily lows attained at the prescribed closes and their daily highs below the permitted ceilings. Within-step interval ratios enclose both $|u|/x$ and $9|u|/y$ and give the gross-rate bound in (7.19). The same enclosures, with outward-rounded time sums, bound individual gross share sales and executed-dollar turnover, as claimed.

Finally, the terminal mean loss is $-\log[(0.9409)^9 y(40)]/10$. The assertion that it exceeds 0.020 is checked *without* trusting a floating-point logarithm: finite rational Taylor series provide opposite bounds on $e^{0.2}$ and verify $e^{0.2}(0.9409)^9 \sup y(40) < 1$. A second rational check at $e^{0.21}$ confirms the mean loss remains below 0.021. These assertions certify feasibility of the stated piecewise-linear path and its compensator ODE.

References

[1] A. Alfonsi, F. Klöck, and A. Schied, *Multivariate transient price impact and matrix-valued positive definite functions*, Mathematics of Operations Research, 41 (2016), pp. 914–934.

[2] A. Alfonsi and A. Schied, *Optimal trade execution and absence of price manipulations in limit order book models*, SIAM Journal on Financial Mathematics, 1 (2010), pp. 490–522.

[3] A. Alfonsi, A. Schied, and A. Slynko, *Order book resilience, price manipulation, and the positive portfolio problem*, SIAM Journal on Financial Mathematics, 3 (2012), pp. 511–533.

[4] G. Cespa and T. Foucault, *Illiquidity contagion and liquidity crashes*, The Review of Financial Studies, 27 (2014), pp. 1615–1660.

[5] T. Chen, Z. Gao, J. He, W. Jiang, and W. Xiong, *Daily price limits and destructive market behavior*, Journal of Econometrics, 208 (2019), pp. 249–264.

[6] R. Cont and L. Wagalath, *Running for the exit: Distressed selling and endogenous correlation in financial markets*, Mathematical Finance, 23 (2013), pp. 718–741.

[7] F. Cordoni and F. Lillo, *Instabilities in multi-asset and multi-agent market impact games*, Annals of Operations Research, 336 (2024), pp. 505–539.

[8] J. Gatheral, *No-dynamic-arbitrage and market impact*, Quantitative Finance, 10 (2010), pp. 749–759.

[9] G. Gennotte and H. Leland, *Market liquidity, hedging, and crashes*, American Economic Review, 80 (1990), pp. 999–1021.

[10] J. Huang and J. Wang, *Liquidity and market crashes*, The Review of Financial Studies, 22 (2009), pp. 2607–2643.

[11] G. Huberman and W. Stanzl, *Price manipulation and quasi-arbitrage*, Econometrica, 72 (2004), pp. 1247–1275.

[12] Japan Exchange Group, *Daily price limits: Trading rules of domestic stocks*. https://www.jpx.co.jp/english/equities/trading/domestic/06.html, 2026. Accessed October 2026.

[13] K. A. Kim and S. G. Rhee, *Price limit performance: Evidence from the Tokyo Stock Exchange*, The Journal of Finance, 52 (1997), pp. 885–901.

[14] A. S. Kyle, *Continuous auctions and insider trading*, Econometrica, 53 (1985), pp. 1315–1336.

[15] A. S. Kyle and A. A. Obizhaeva, *Large bets and stock market crashes*, Review of Finance, 27 (2023), pp. 2163–2203.

[16] A. A. Obizhaeva and J. Wang, *Optimal trading strategy and supply/demand dynamics*, Journal of Financial Markets, 16 (2013), pp. 1–32.

[17] D. Romer, *Rational asset-price movements without news*, American Economic Review, 83 (1993), pp. 1112–1130.

[18] Y. Yang, *Finite horizon optimal execution with bounded rate of transaction*, Stochastic Models, 35 (2019), pp. 469–495.